

Optimization Modulo Theory: A Tutorial Using Z3 and Practical Case Studies

Mădălina Erașcu

West University of Timișoara, Romania

madalina.erascu@e-uvt.ro

September 16th, 2024

This work was supported by a grant of the Romanian National Authority for Scientific Research and Innovation, CNCS/CCCDI-UEFISCDI, project number PN-III-P1-1.1-TE-2021-0676.

Outline

Motivation

Part 1: Optimization Modulo Theory - background and examples

Part 2: Optimization Modulo Theory - Case Study

Problem Specification

Problem Formalization

Model-driven approach: Formulation of the Satisfiability/Optimization Modulo Theory Problem

Data-driven approach: Graph Neural Network Formulation

Solution

Dataset generation

Training a GNN model for edge classification

Integrated GNN and Exact Techniques: Experimental Results

Future Work

Acknowledgements

- ▶ Material resources: Erika Abraham (RWTH Aachen),
<https://microsoft.github.io/z3guide/>, <https://github.com/Z3Prover/z3>.

Z3 solver online to be used during the tutorial



Acknowledgements

- ▶ Material resources: Erika Abraham (RWTH Aachen),
<https://microsoft.github.io/z3guide/>, <https://github.com/Z3Prover/z3>.
- ▶ Parts of this work is joint with: Flavia Micota, Daniela Zaharie, Eduard Laitin.

Acknowledgements

- ▶ Material resources: Erika Abraham (RWTH Aachen),
<https://microsoft.github.io/z3guide/>, <https://github.com/Z3Prover/z3>.
- ▶ Parts of this work is joint with: Flavia Micota, Daniela Zaharie, Eduard Laitin.
- ▶ Parts of this work was inspired by or polished after teaching satisfiability checking topics to Master students or supervising Bachelor/Master theses.

Acknowledgements

- ▶ Material resources: Erika Abraham (RWTH Aachen),
<https://microsoft.github.io/z3guide/>, <https://github.com/Z3Prover/z3>.
- ▶ Parts of this work is joint with: Flavia Micota, Daniela Zaharie, Eduard Laitin.
- ▶ Parts of this work was inspired by or polished after teaching satisfiability checking topics to Master students or supervising Bachelor/Master theses.

Files used in this presentation



Contents

Motivation

Part 1: Optimization Modulo Theory - background and examples

Part 2: Optimization Modulo Theory - Case Study

Problem Specification

Problem Formalization

Model-driven approach: Formulation of the Satisfiability/Optimization Modulo Theory Problem

Data-driven approach: Graph Neural Network Formulation

Solution

Dataset generation

Training a GNN model for edge classification

Integrated GNN and Exact Techniques: Experimental Results

Future Work

Motivation

Constrained optimization problems have applications in engineering, economics, computer science, just to name a few.

Motivation

Constrained optimization problems have applications in engineering, economics, computer science, just to name a few.

General form:

$$\begin{array}{ll}\min & f_1(\mathbf{x}), f_2(\mathbf{x}), \dots, f_p(\mathbf{x}) \quad \forall p \geq 1 \\ \text{subject to} & g_i(\mathbf{x}) = c_i \quad \forall i = \overline{1, n} \\ & h_j(\mathbf{x}) \geq d_j \quad \forall j = \overline{1, m}\end{array}$$

Motivation

Constrained optimization problems have applications in engineering, economics, computer science, just to name a few.

General form:

$$\begin{array}{ll}\min & f_1(\mathbf{x}), f_2(\mathbf{x}), \dots, f_p(\mathbf{x}) \quad \forall p \geq 1 \\ \text{subject to} & g_i(\mathbf{x}) = c_i \quad \forall i = \overline{1, n} \\ & h_j(\mathbf{x}) \geq d_j \quad \forall j = \overline{1, m}\end{array}$$

Solution approaches:

Motivation

Constrained optimization problems have applications in engineering, economics, computer science, just to name a few.

General form:

$$\begin{array}{ll}\min & f_1(\mathbf{x}), f_2(\mathbf{x}), \dots, f_p(\mathbf{x}) \quad \forall p \geq 1 \\ \text{subject to} & g_i(\mathbf{x}) = c_i \quad \forall i = \overline{1, n} \\ & h_j(\mathbf{x}) \geq d_j \quad \forall j = \overline{1, m}\end{array}$$

Solution approaches:

1. Exact methods

- ▶ Constrained Programming (CP)
 - ▶ Modelling languages can be used, e.g. [MiniZinc](#) [15]
 - ▶ [Google OR-Tools](#) [16], [Gecode](#) [17], [Chuffed](#) [6]

Motivation

Constrained optimization problems have applications in engineering, economics, computer science, just to name a few.

General form:

$$\begin{array}{ll} \min & f_1(\mathbf{x}), f_2(\mathbf{x}), \dots, f_p(\mathbf{x}) \quad \forall p \geq 1 \\ \text{subject to} & g_i(\mathbf{x}) = c_i \quad \forall i = \overline{1, n} \\ & h_j(\mathbf{x}) \geq d_j \quad \forall j = \overline{1, m} \end{array}$$

Solution approaches:

1. Exact methods

- ▶ **Constrained Programming (CP)**
 - ▶ Modelling languages can be used, e.g. [MiniZinc](#) [15]
 - ▶ [Google OR-Tools](#) [16], [Gecode](#) [17], [Chuffed](#) [6]
- ▶ **Mathematical Programming (MP)**
 - ▶ [CPLEX](#) [1]

Motivation

Constrained optimization problems have applications in engineering, economics, computer science, just to name a few.

General form:

$$\begin{array}{ll} \min & f_1(\mathbf{x}), f_2(\mathbf{x}), \dots, f_p(\mathbf{x}) \quad \forall p \geq 1 \\ \text{subject to} & g_i(\mathbf{x}) = c_i \quad \forall i = \overline{1, n} \\ & h_j(\mathbf{x}) \geq d_j \quad \forall j = \overline{1, m} \end{array}$$

Solution approaches:

1. Exact methods

- ▶ Constrained Programming (CP)
 - ▶ Modelling languages can be used, e.g. [MiniZinc](#) [15]
 - ▶ [Google OR-Tools](#) [16], [Gecode](#) [17], [Chuffed](#) [6]
- ▶ Mathematical Programming (MP)
 - ▶ [CPLEX](#) [1]
- ▶ Satisfiability Modulo Theory (SMT)
 - ▶ [Z3](#) [7], [cvc5](#) [3], [Yices2](#) [8], [OpenSMT](#) [5].
- ▶ Advantage: provides an optimal solution

Motivation

Constrained optimization problems have applications in engineering, economics, computer science, just to name a few.

General form:

$$\begin{array}{ll}\min & f_1(\mathbf{x}), f_2(\mathbf{x}), \dots, f_p(\mathbf{x}) \quad \forall p \geq 1 \\ \text{subject to} & g_i(\mathbf{x}) = c_i \quad \forall i = \overline{1, n} \\ & h_j(\mathbf{x}) \geq d_j \quad \forall j = \overline{1, m}\end{array}$$

Solution approaches:

1. Exact methods

- ▶ Constrained Programming (CP)
 - ▶ Modelling languages can be used, e.g. [MiniZinc](#) [15]
 - ▶ [Google OR-Tools](#) [16], [Gecode](#) [17], [Chuffed](#) [6]
- ▶ Mathematical Programming (MP)
 - ▶ [CPLEX](#) [1]
- ▶ Satisfiability Modulo Theory (SMT)
 - ▶ [Z3](#) [7], [cvc5](#) [3], [Yices2](#) [8], [OpenSMT](#) [5].
- ▶ Advantage: provides an optimal solution
- ▶ Drawback: significant computational time for large problems

Motivation

Constrained optimization problems have applications in engineering, economics, computer science, just to name a few.

General form:

$$\begin{array}{ll} \min & f_1(\mathbf{x}), f_2(\mathbf{x}), \dots, f_p(\mathbf{x}) \quad \forall p \geq 1 \\ \text{subject to} & g_i(\mathbf{x}) = c_i \quad \forall i = \overline{1, n} \\ & h_j(\mathbf{x}) \geq d_j \quad \forall j = \overline{1, m} \end{array}$$

Solution approaches:

1. Exact methods

- ▶ Constrained Programming (CP)
 - ▶ Modelling languages can be used, e.g. MiniZinc [15]
 - ▶ Google OR-Tools [16], Gecode [17], Chuffed [6]
- ▶ Mathematical Programming (MP)
 - ▶ CPLEX [1]
- ▶ Satisfiability Modulo Theory (SMT)
 - ▶ Z3 [7], cvc5 [3], Yices2 [8], OpenSMT [5].
- ▶ Advantage: provides an optimal solution
- ▶ Drawback: significant computational time for large problems

2. Approximate methods

Motivation

Constrained optimization problems have applications in engineering, economics, computer science, just to name a few.

General form:

$$\begin{array}{ll} \min & f_1(\mathbf{x}), f_2(\mathbf{x}), \dots, f_p(\mathbf{x}) \quad \forall p \geq 1 \\ \text{subject to} & g_i(\mathbf{x}) = c_i \quad \forall i = \overline{1, n} \\ & h_j(\mathbf{x}) \geq d_j \quad \forall j = \overline{1, m} \end{array}$$

Solution approaches:

1. Exact methods

- ▶ Constrained Programming (CP)
 - ▶ Modelling languages can be used, e.g. MiniZinc [15]
 - ▶ Google OR-Tools [16], Gecode [17], Chuffed [6]
- ▶ Mathematical Programming (MP)
 - ▶ CPLEX [1]
- ▶ Satisfiability Modulo Theory (SMT)
 - ▶ Z3 [7], cvc5 [3], Yices2 [8], OpenSMT [5].
- ▶ Advantage: provides an optimal solution
- ▶ Drawback: significant computational time for large problems

2. Approximate methods

- ▶ heuristics and metaheuristics

Motivation

Constrained optimization problems have applications in engineering, economics, computer science, just to name a few.

General form:

$$\begin{array}{ll} \min & f_1(\mathbf{x}), f_2(\mathbf{x}), \dots, f_p(\mathbf{x}) \quad \forall p \geq 1 \\ \text{subject to} & g_i(\mathbf{x}) = c_i \quad \forall i = \overline{1, n} \\ & h_j(\mathbf{x}) \geq d_j \quad \forall j = \overline{1, m} \end{array}$$

Solution approaches:

1. Exact methods

- ▶ Constrained Programming (CP)
 - ▶ Modelling languages can be used, e.g. [MiniZinc](#) [15]
 - ▶ [Google OR-Tools](#) [16], [Gecode](#) [17], [Chuffed](#) [6]
- ▶ Mathematical Programming (MP)
 - ▶ [CPLEX](#) [1]
- ▶ Satisfiability Modulo Theory (SMT)
 - ▶ [Z3](#) [7], [cvc5](#) [3], [Yices2](#) [8], [OpenSMT](#) [5].
- ▶ Advantage: provides an optimal solution
- ▶ Drawback: significant computational time for large problems

2. Approximate methods

- ▶ heuristics and metaheuristics
- ▶ Advantage: faster

Motivation

Constrained optimization problems have applications in engineering, economics, computer science, just to name a few.

General form:

$$\begin{array}{ll} \min & f_1(\mathbf{x}), f_2(\mathbf{x}), \dots, f_p(\mathbf{x}) \quad \forall p \geq 1 \\ \text{subject to} & g_i(\mathbf{x}) = c_i \quad \forall i = \overline{1, n} \\ & h_j(\mathbf{x}) \geq d_j \quad \forall j = \overline{1, m} \end{array}$$

Solution approaches:

1. Exact methods

- ▶ Constrained Programming (CP)
 - ▶ Modelling languages can be used, e.g. MiniZinc [15]
 - ▶ Google OR-Tools [16], Gecode [17], Chuffed [6]
- ▶ Mathematical Programming (MP)
 - ▶ CPLEX [1]
- ▶ Satisfiability Modulo Theory (SMT)
 - ▶ Z3 [7], cvc5 [3], Yices2 [8], OpenSMT [5].
- ▶ Advantage: provides an optimal solution
- ▶ Drawback: significant computational time for large problems

2. Approximate methods

- ▶ heuristics and metaheuristics
- ▶ Advantage: faster
- ▶ Drawback: provides a (sub)optimal solution

Contents

Motivation

Part 1: Optimization Modulo Theory - background and examples

Part 2: Optimization Modulo Theory - Case Study

Problem Specification

Problem Formalization

Model-driven approach: Formulation of the Satisfiability/Optimization Modulo Theory Problem

Data-driven approach: Graph Neural Network Formulation

Solution

Dataset generation

Training a GNN model for edge classification

Integrated GNN and Exact Techniques: Experimental Results

Future Work

Satisfiability Modulo Theory (SMT)

- Satisfiability is the problem of determining if a formula has a model.

Satisfiability Modulo Theory (SMT)

- ▶ **Satisfiability** is the problem of determining if a formula has a model.
- ▶ In the **propositional case**, a model is a truth assignment to the Boolean variables.

Satisfiability Modulo Theory (SMT)

- ▶ Satisfiability is the problem of determining if a formula has a model.
- ▶ In the propositional case, a model is a truth assignment to the Boolean variables.
- ▶ In the first-order (FO) case, a model assigns values from a domain to variables and interpretations over the domain to the function and predicate symbols.

Satisfiability Modulo Theory (SMT)

- ▶ Satisfiability is the problem of determining if a formula has a model.
- ▶ In the propositional case, a model is a truth assignment to the Boolean variables.
- ▶ In the first-order (FO) case, a model assigns values from a domain to variables and interpretations over the domain to the function and predicate symbols.
- ▶ Automated reasoning failure: proof-search procedures for full FO logic \rightsquigarrow is FO logic the best compromise between expressivity and efficiency?

Satisfiability Modulo Theory (SMT)

- ▶ Satisfiability is the problem of determining if a formula has a model.
- ▶ In the propositional case, a model is a truth assignment to the Boolean variables.
- ▶ In the first-order (FO) case, a model assigns values from a domain to variables and interpretations over the domain to the function and predicate symbols.
- ▶ Automated reasoning failure: proof-search procedures for full FO logic \rightsquigarrow is FO logic the best compromise between expressivity and efficiency?
- ▶ Gain efficiency by:

Satisfiability Modulo Theory (SMT)

- ▶ Satisfiability is the problem of determining if a formula has a model.
- ▶ In the propositional case, a model is a truth assignment to the Boolean variables.
- ▶ In the first-order (FO) case, a model assigns values from a domain to variables and interpretations over the domain to the function and predicate symbols.
- ▶ Automated reasoning failure: proof-search procedures for full FO logic \rightsquigarrow is FO logic the best compromise between expressivity and efficiency?
- ▶ Gain efficiency by:
 - ▶ addressing only (expressive enough) decidable fragments of a certain logic.

Satisfiability Modulo Theory (SMT)

- ▶ Satisfiability is the problem of determining if a formula has a model.
- ▶ In the propositional case, a model is a truth assignment to the Boolean variables.
- ▶ In the first-order (FO) case, a model assigns values from a domain to variables and interpretations over the domain to the function and predicate symbols.
- ▶ Automated reasoning failure: proof-search procedures for full FO logic \rightsquigarrow is FO logic the best compromise between expressivity and efficiency?
- ▶ Gain efficiency by:
 - ▶ addressing only (expressive enough) decidable fragments of a certain logic.
 - ▶ incorporate domain-specific reasoning, e.g: arithmetic, equality, data structures (arrays, lists, stacks, ...) and valid combinations

Satisfiability (SMT) and Optimization Modulo Theory (OMT) (cont'd)

- **SAT**: uses propositional logic as the formalization language

Satisfiability (SMT) and Optimization Modulo Theory (OMT) (cont'd)

- ▶ **SAT**: uses propositional logic as the formalization language
- ▶ **SMT**: propositional logic + domain-specific reasoning

Satisfiability (SMT) and Optimization Modulo Theory (OMT) (cont'd)

- ▶ **SAT**: uses propositional logic as the formalization language
 - ▶ +: high degree of efficiency
- ▶ **SMT**: propositional logic + domain-specific reasoning

Satisfiability (SMT) and Optimization Modulo Theory (OMT) (cont'd)

- ▶ **SAT**: uses propositional logic as the formalization language
 - ▶ +: high degree of efficiency
 - ▶ -: expressive but complex encodings
- ▶ **SMT**: propositional logic + domain-specific reasoning

Satisfiability (SMT) and Optimization Modulo Theory (OMT) (cont'd)

- ▶ **SAT**: uses propositional logic as the formalization language
 - ▶ +: high degree of efficiency
 - ▶ -: expressive but complex encodings
- ▶ **SMT**: propositional logic + domain-specific reasoning
 - ▶ + better expressivity

Satisfiability (SMT) and Optimization Modulo Theory (OMT) (cont'd)

- ▶ **SAT**: uses propositional logic as the formalization language
 - ▶ +: high degree of efficiency
 - ▶ -: expressive but complex encodings
- ▶ **SMT**: propositional logic + domain-specific reasoning
 - ▶ + better expressivity
 - ▶ -: certain (but acceptable) loss of efficiency

Satisfiability (SMT) and Optimization Modulo Theory (OMT) (cont'd)

- ▶ **SAT**: uses propositional logic as the formalization language
 - ▶ +: high degree of efficiency
 - ▶ -: expressive but complex encodings
- ▶ **SMT**: propositional logic + domain-specific reasoning
 - ▶ + better expressivity
 - ▶ -: certain (but acceptable) loss of efficiency
- ▶ SAT competition: <https://satcompetition.github.io/>

Satisfiability (SMT) and Optimization Modulo Theory (OMT) (cont'd)

- ▶ **SAT**: uses propositional logic as the formalization language
 - ▶ +: high degree of efficiency
 - ▶ -: expressive but complex encodings
- ▶ **SMT**: propositional logic + domain-specific reasoning
 - ▶ + better expressivity
 - ▶ -: certain (but acceptable) loss of efficiency
- ▶ SAT competition: <https://satcompetition.github.io/>
- ▶ SMT competition: <https://smt-comp.github.io/>

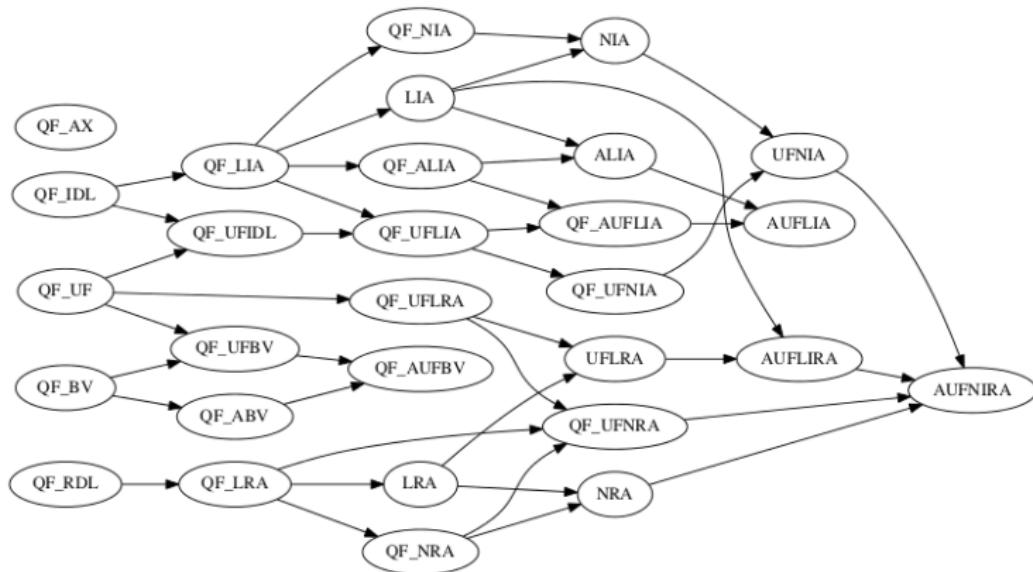
Satisfiability (SMT) and Optimization Modulo Theory (OMT) (cont'd)

- ▶ **SAT**: uses propositional logic as the formalization language
 - ▶ +: high degree of efficiency
 - ▶ -: expressive but complex encodings
- ▶ **SMT**: propositional logic + domain-specific reasoning
 - ▶ + better expressivity
 - ▶ -: certain (but acceptable) loss of efficiency
- ▶ SAT competition: <https://satcompetition.github.io/>
- ▶ SMT competition: <https://smt-comp.github.io/>

Satisfiability (SMT) and Optimization Modulo Theory (OMT) (cont'd)

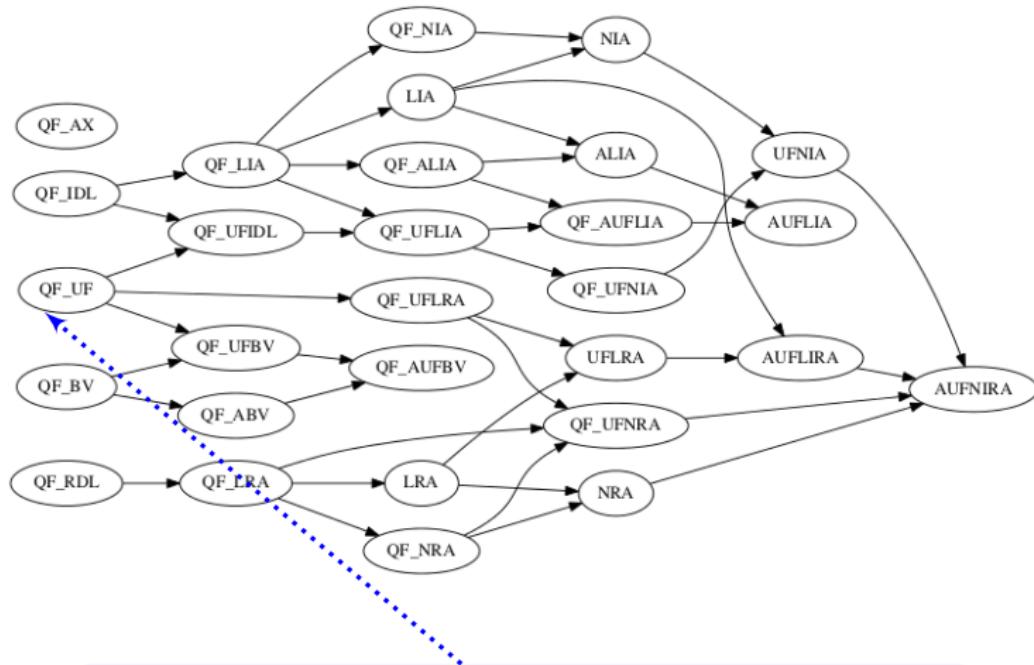
- ▶ **SAT**: uses propositional logic as the formalization language
 - ▶ +: high degree of efficiency
 - ▶ -: expressive but complex encodings
- ▶ **SMT**: propositional logic + domain-specific reasoning
 - ▶ + better expressivity
 - ▶ -: certain (but acceptable) loss of efficiency
- ▶ SAT competition: <https://satcompetition.github.io/>
- ▶ SMT competition: <https://smt-comp.github.io/>
- ▶ Some SMT solvers offer optimization features \rightsquigarrow optimization modulo theory (OMT): Z3 [4], OptiMathSAT [18]; Symba [13], HAZEL [14], MAXHS-MSAT [10], PULI [11], CEGIO [2], BCLET [12].

SMT Theories



Source: <http://smtlib.cs.uiowa.edu/logics.shtml>

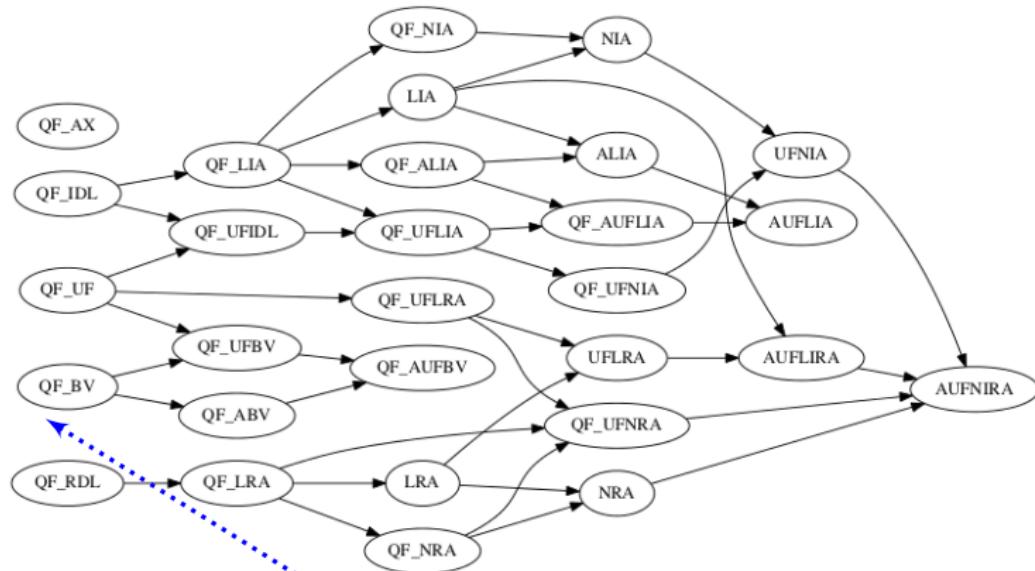
SMT Theories



Quantifier-free equality logic with uninterpreted functions
 $(a = c \wedge b = d) \rightarrow f(a, b) = f(c, d)$

Source: <http://smtlib.cs.uiowa.edu/logics.shtml>

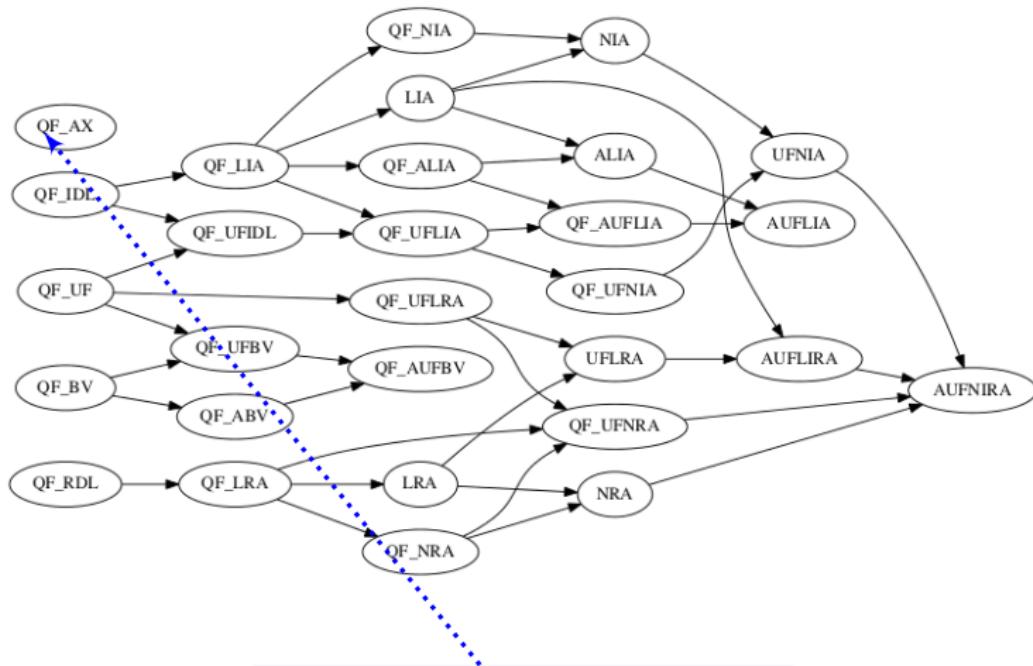
SMT Theories



Quantifier-free bit-vector arithmetic
 $a + b \geq 0 \wedge (a|b) \leq (a\&b)$

Source: <http://smtlib.cs.uiowa.edu/logics.shtml>

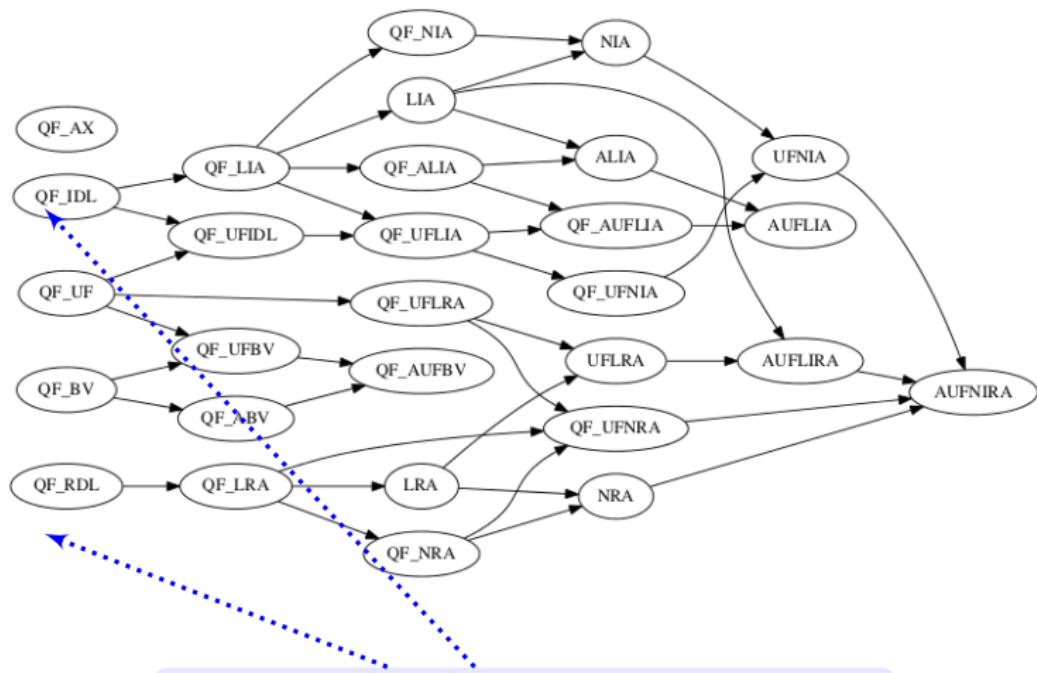
SMT Theories



Quantifier-free array theory
 $i = j \rightarrow \text{read}(\text{write}(a, i, v), j) = v$

Source: <http://smtlib.cs.uiowa.edu/logics.shtml>

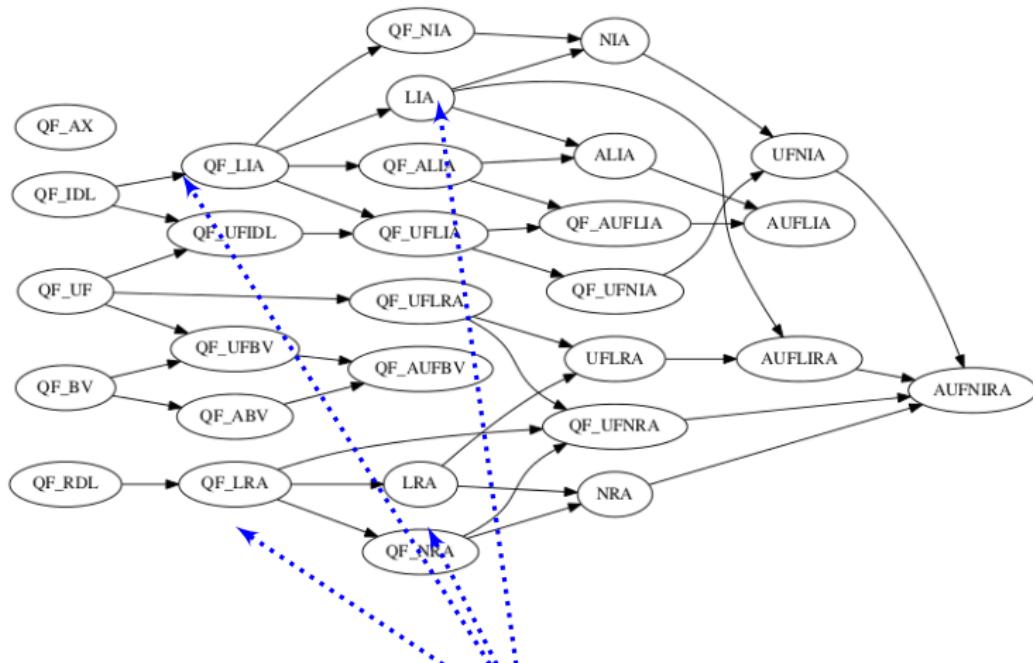
SMT Theories



Quantifier-free integer/rational difference logic
 $x - y \geq 0 \vee x - z < 0$

Source: <http://smtlib.cs.uiowa.edu/logics.shtml>

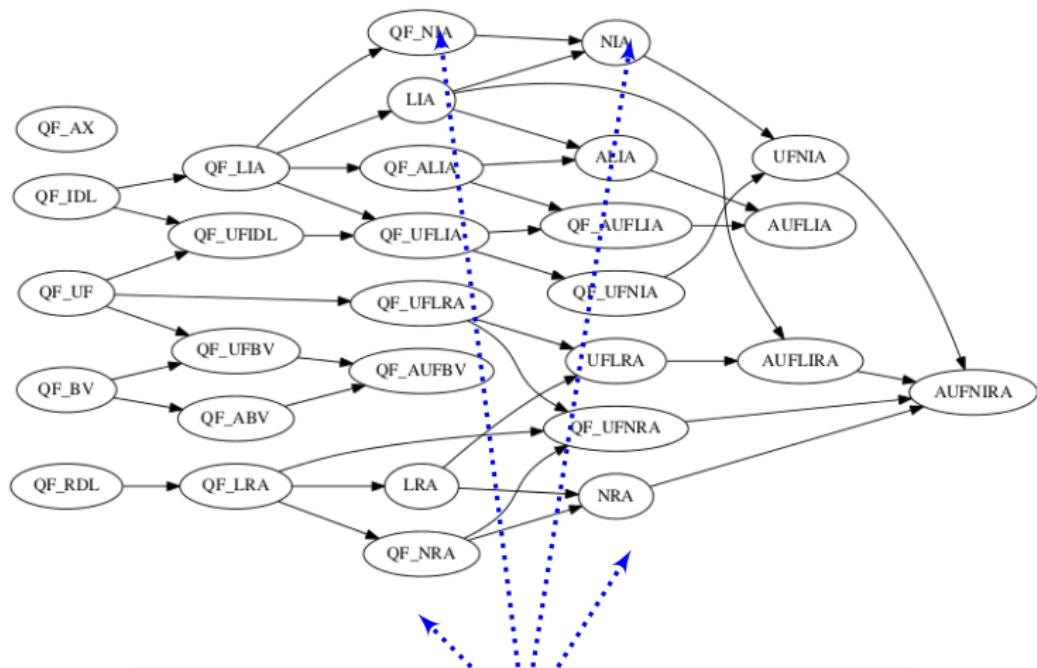
SMT Theories



(Quantifier-free) real/integer linear arithmetic
$$4x + 7y = 8 \wedge (y = 0 \vee x > y)$$

Source: <http://smtlib.cs.uiowa.edu/logics.shtml>

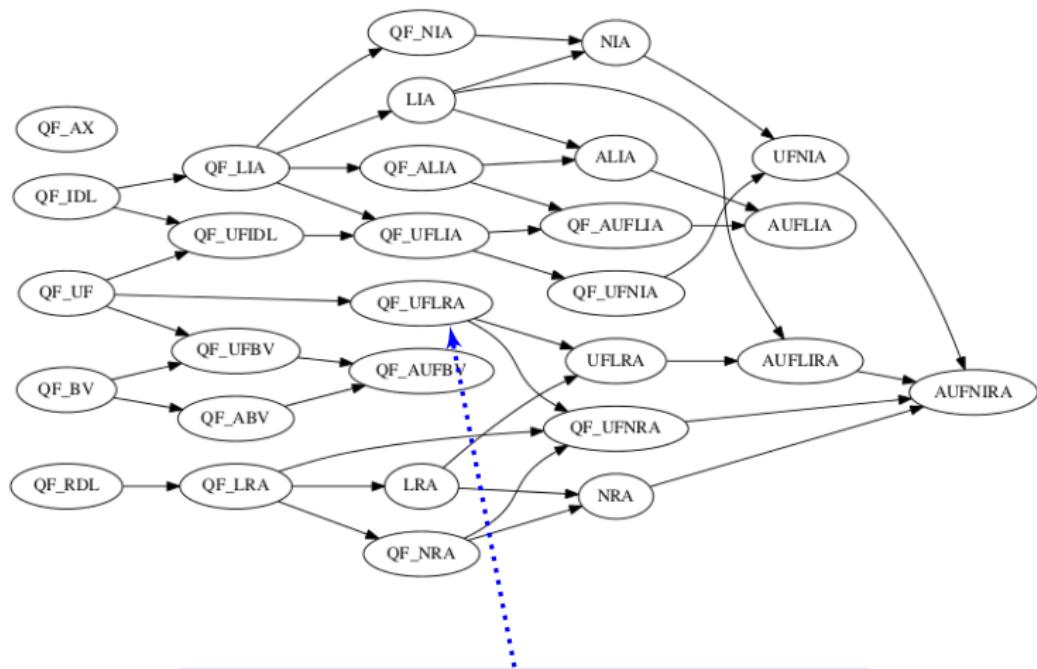
SMT Theories



(Quantifier-free) real/integer non-linear arithmetic
 $x^2 + 2xy + y^2 > 0 \vee (x \geq 1 \wedge xz + yz^2 = 0)$

Source: <http://smtlib.cs.uiowa.edu/logics.shtml>

SMT Theories



Combined theories

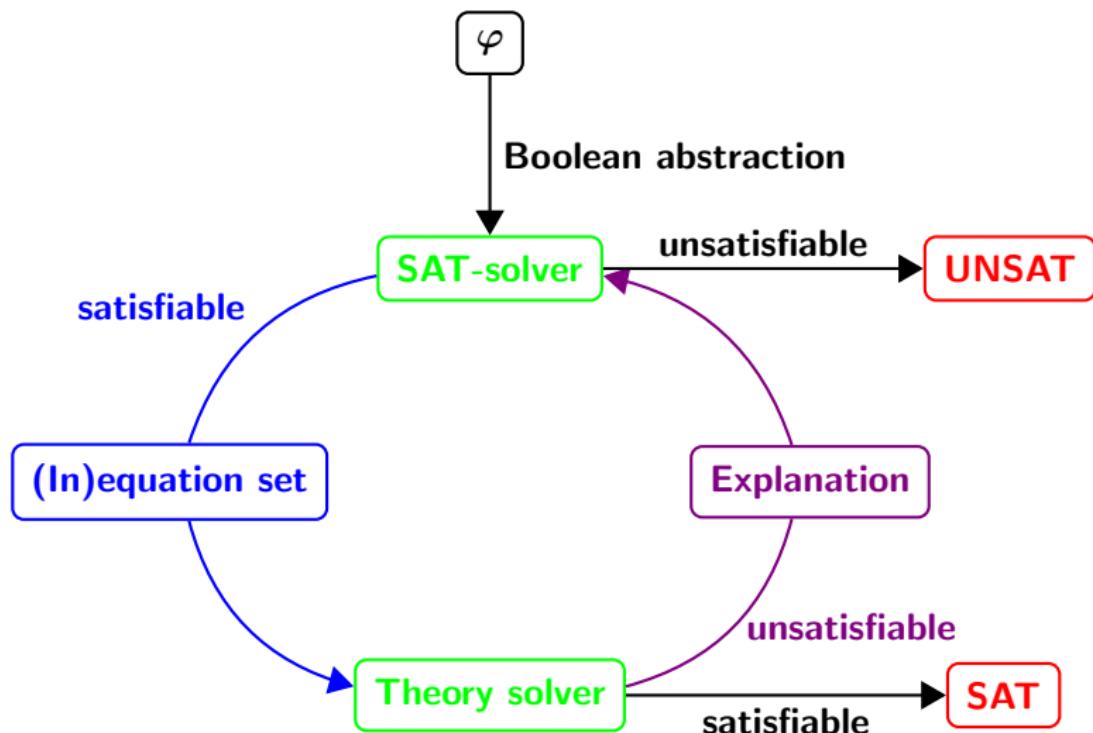
Source: <http://smtlib.cs.uiowa.edu/logics.shtml>

How an extension to SMT solving looks like?

There are basically two different approaches:

- ▶ **Eager SMT solving** transforms logical formulas over some theories into satisfiability-equivalent propositional logic formulas and applies SAT solving. (“Eager” means theory first)
- ▶ **Lazy SMT solving** uses a SAT solver to find solutions for the Boolean skeleton of the formula, and a theory solver to check satisfiability in the underlying theory. (“Lazy” means theory later)

Lazy SMT solving



Running Example

Assume that we have three virtual machines (VMs) which require 100, 50 and 15 GB hard disk respectively. There are three servers with capabilities 100, 75 and 200 GB in that order. Find out a way to place VMs into servers in order to:

- ▶ Minimize the number of servers used.
- ▶ Minimize the operation cost (the servers have fixed daily costs 10, 5 and 20 USD respectively.)

Running Example

Assume that we have three virtual machines (VMs) which require 100, 50 and 15 GB hard disk respectively. There are three servers with capabilities 100, 75 and 200 GB in that order. Find out a way to place VMs into servers in order to:

- ▶ Minimize the number of servers used.
- ▶ Minimize the operation cost (the servers have fixed daily costs 10, 5 and 20 USD respectively.)

Formalization. Let x_{ij} denote that VM i is placed on the server j and y_j denote that server j is in use.

Running Example (cont'd)

Assume that we have three virtual machines (VMs) which require 100, 50 and 15 GB hard disk respectively. There are three servers with capabilities 100, 75 and 200 GB in that order. Find out a way to place VMs into servers in order to:

- ▶ Minimize the number of servers used.
- ▶ Minimize the operation cost (the servers have fixed daily costs 10, 5 and 20 USD respectively.)

Solution. Choosing the suitable underlying theory is determined by the principles of the formalization: $x_{ij}, y_j \in \{0, 1\}$

- ▶ linear constraints with integer variables with 0,1 restriction
- ▶ linear constraints with real variables with 0,1 restriction
- ▶ linear constraints boolean variables

Running Example (cont'd)

Assume that we have three virtual machines (VMs) which require 100, 50 and 15 GB hard disk respectively. There are three servers with capabilities 100, 75 and 200 GB in that order. Find out a way to place VMs into servers in order to:

- ▶ Minimize the number of servers used.
- ▶ Minimize the operation cost (the servers have fixed daily costs 10, 5 and 20 USD respectively.)

Solution. Writing the constraints keeping in mind the underlying theory.

Assume the case: linear constraints with integer variables with 0,1 restriction

- ▶ Implicit constraints

Running Example (cont'd)

Assume that we have three virtual machines (VMs) which require 100, 50 and 15 GB hard disk respectively. There are three servers with capabilities 100, 75 and 200 GB in that order. Find out a way to place VMs into servers in order to:

- ▶ Minimize the number of servers used.
- ▶ Minimize the operation cost (the servers have fixed daily costs 10, 5 and 20 USD respectively.)

Solution. Writing the constraints keeping in mind the underlying theory.

Assume the case: linear constraints with integer variables with 0,1 restriction

- ▶ Implicit constraints
 - ▶ Variables are integers:

$$x_{ij}, y_j \in \mathbb{Z}, \quad \forall i, j = \overline{1, 3}$$

Running Example (cont'd)

Assume that we have three virtual machines (VMs) which require 100, 50 and 15 GB hard disk respectively. There are three servers with capabilities 100, 75 and 200 GB in that order. Find out a way to place VMs into servers in order to:

- ▶ Minimize the number of servers used.
- ▶ Minimize the operation cost (the servers have fixed daily costs 10, 5 and 20 USD respectively.)

Solution. Writing the constraints keeping in mind the underlying theory.

Assume the case: linear constraints with integer variables with 0,1 restriction

- ▶ Implicit constraints

- ▶ Variables are integers:

$$x_{ij}, y_j \in \mathbb{Z}, \quad \forall i, j = \overline{1, 3}$$

- ▶ Variables have only 0,1 value:

$$x_{ij} = 0 \vee x_{ij} = 1, \quad \forall i, j = \overline{1, 3}$$

Running Example (cont'd)

Assume that we have three virtual machines (VMs) which require 100, 50 and 15 GB hard disk respectively. There are three servers with capabilities 100, 75 and 200 GB in that order. Find out a way to place VMs into servers in order to:

- ▶ Minimize the number of servers used.
- ▶ Minimize the operation cost (the servers have fixed daily costs 10, 5 and 20 USD respectively.)

Solution. Writing the constraints keeping in mind the underlying theory.

Assume the case: linear constraints with integer variables with 0,1 restriction

- ▶ Implicit constraints

- ▶ Variables are integers:

$$x_{ij}, y_j \in \mathbb{Z}, \quad \forall i, j = \overline{1, 3}$$

- ▶ Variables have only 0,1 value:

$$x_{ij} = 0 \vee x_{ij} = 1, \quad \forall i, j = \overline{1, 3}$$

- ▶ A VM is on exactly one server:

$$x_{i1} + x_{i2} + x_{i3} = 1, \quad \forall i = \overline{1, 3}$$

Running Example (cont'd)

Assume that we have three virtual machines (VMs) which require 100, 50 and 15 GB hard disk respectively. There are three servers with capabilities 100, 75 and 200 GB in that order. Find out a way to place VMs into servers in order to:

- ▶ Minimize the number of servers used.
- ▶ Minimize the operation cost (the servers have fixed daily costs 10, 5 and 20 USD respectively.)

Solution. Writing the constraints keeping in mind the underlying theory.

Assume the case: linear constraints with integer variables with 0,1 restriction

- ▶ **Implicit constraints**

- ▶ Variables are integers:

$$x_{ij}, y_j \in \mathbb{Z}, \quad \forall i, j = \overline{1, 3}$$

- ▶ Variables have only 0,1 value:

$$x_{ij} = 0 \vee x_{ij} = 1, \quad \forall i, j = \overline{1, 3}$$

- ▶ A VM is on exactly one server:

$$x_{i1} + x_{i2} + x_{i3} = 1, \quad \forall i = \overline{1, 3}$$

- ▶ A used server has at least a VM on it:

$$(y_j \geq x_{1j}) \wedge (y_j \geq x_{2j}) \wedge (y_j \geq x_{3j}), \quad j = \overline{1, 3}$$

Running Example (cont'd)

Assume that we have three virtual machines (VMs) which require 100, 50 and 15 GB hard disk respectively. There are three servers with capabilities 100, 75 and 200 GB in that order. Find out a way to place VMs into servers in order to:

- ▶ Minimize the number of servers used.
- ▶ Minimize the operation cost (the servers have fixed daily costs 10, 5 and 20 USD respectively.)

Solution. Writing the constraints keeping in mind the underlying theory.

Assume the case: linear constraints with integer variables with 0,1 restriction

- ▶ Explicit constraints

Running Example (cont'd)

Assume that we have three virtual machines (VMs) which require 100, 50 and 15 GB hard disk respectively. There are three servers with capabilities 100, 75 and 200 GB in that order. Find out a way to place VMs into servers in order to:

- ▶ Minimize the number of servers used.
- ▶ Minimize the operation cost (the servers have fixed daily costs 10, 5 and 20 USD respectively.)

Solution. Writing the constraints keeping in mind the underlying theory.

Assume the case: linear constraints with integer variables with 0,1 restriction

- ▶ Explicit constraints

- ▶ Capacity constraints:

$$100x_{11} + 50x_{21} + 15x_{31} \leq 100y_1$$

$$100x_{12} + 50x_{22} + 15x_{32} \leq 75y_2$$

$$100x_{13} + 50x_{23} + 15x_{33} \leq 200y_3$$

Running Example (cont'd)

Assume that we have three virtual machines (VMs) which require 100, 50 and 15 GB hard disk respectively. There are three servers with capabilities 100, 75 and 200 GB in that order. Find out a way to place VMs into servers in order to:

- ▶ Minimize the number of servers used.
- ▶ Minimize the operation cost (the servers have fixed daily costs 10, 5 and 20 USD respectively.)

Solution. Writing the constraints keeping in mind the underlying theory.

Assume the case: linear constraints with integer variables with 0,1 restriction

- ▶ Explicit constraints

- ▶ Capacity constraints:

$$100x_{11} + 50x_{21} + 15x_{31} \leq 100y_1$$

$$100x_{12} + 50x_{22} + 15x_{32} \leq 75y_2$$

$$100x_{13} + 50x_{23} + 15x_{33} \leq 200y_3$$

- ▶ Optimization functions

Running Example (cont'd)

Assume that we have three virtual machines (VMs) which require 100, 50 and 15 GB hard disk respectively. There are three servers with capabilities 100, 75 and 200 GB in that order. Find out a way to place VMs into servers in order to:

- ▶ Minimize the number of servers used.
- ▶ Minimize the operation cost (the servers have fixed daily costs 10, 5 and 20 USD respectively.)

Solution. Writing the constraints keeping in mind the underlying theory.

Assume the case: linear constraints with integer variables with 0,1 restriction

- ▶ **Explicit constraints**

- ▶ Capacity constraints:

$$100x_{11} + 50x_{21} + 15x_{31} \leq 100y_1$$

$$100x_{12} + 50x_{22} + 15x_{32} \leq 75y_2$$

$$100x_{13} + 50x_{23} + 15x_{33} \leq 200y_3$$

- ▶ **Optimization functions**

- ▶ $10y_1 + 5y_2 + 20y_3$

- ▶ $y_1 + y_2 + y_3$

Running Example (cont'd)

Assume that we have three virtual machines (VMs) which require 100, 50 and 15 GB hard disk respectively. There are three servers with capabilities 100, 75 and 200 GB in that order. Find out a way to place VMs into servers in order to:

- ▶ Minimize the number of servers used.
- ▶ Minimize the operation cost (the servers have fixed daily costs 10, 5 and 20 USD respectively.)

Solution approaches.

1. Formalization in [SMT-LIB2 format](#): useful for toy examples, some SMT tools are available online to try their capabilities.

Running Example (cont'd)

Assume that we have three virtual machines (VMs) which require 100, 50 and 15 GB hard disk respectively. There are three servers with capabilities 100, 75 and 200 GB in that order. Find out a way to place VMs into servers in order to:

- ▶ Minimize the number of servers used.
- ▶ Minimize the operation cost (the servers have fixed daily costs 10, 5 and 20 USD respectively.)

Solution approaches.

1. Formalization in [SMT-LIB2 format](#): useful for toy examples, some SMT tools are available online to try their capabilities.
2. Formalization by [programming Z3](#), in particular [Python API](#).

Running Example (cont'd)

Assume that we have three virtual machines (VMs) which require 100, 50 and 15 GB hard disk respectively. There are three servers with capabilities 100, 75 and 200 GB in that order. Find out a way to place VMs into servers in order to:

- ▶ Minimize the number of servers used.
- ▶ Minimize the operation cost (the servers have fixed daily costs 10, 5 and 20 USD respectively.)

Solution approaches.

1. Formalization in [SMT-LIB2 format](#): useful for toy examples, some SMT tools are available online to try their capabilities.
 - ▶ [variant-int.smt2](#)
2. Formalization by [programming Z3](#), in particular [Python API](#).

Running Example (cont'd)

Assume that we have three virtual machines (VMs) which require 100, 50 and 15 GB hard disk respectively. There are three servers with capabilities 100, 75 and 200 GB in that order. Find out a way to place VMs into servers in order to:

- ▶ Minimize the number of servers used.
- ▶ Minimize the operation cost (the servers have fixed daily costs 10, 5 and 20 USD respectively.)

Solution approaches.

1. Formalization in [SMT-LIB2 format](#): useful for toy examples, some SMT tools are available online to try their capabilities.
 - ▶ variant-int.smt2
 - ▶ variant-bool.smt2
2. Formalization by [programming Z3](#), in particular [Python API](#).

Running Example (cont'd)

Assume that we have three virtual machines (VMs) which require 100, 50 and 15 GB hard disk respectively. There are three servers with capabilities 100, 75 and 200 GB in that order. Find out a way to place VMs into servers in order to:

- ▶ Minimize the number of servers used.
- ▶ Minimize the operation cost (the servers have fixed daily costs 10, 5 and 20 USD respectively.)

Solution approaches.

1. Formalization in [SMT-LIB2 format](#): useful for toy examples, some SMT tools are available online to try their capabilities.
 - ▶ variant-int.smt2
 - ▶ variant-bool.smt2
2. Formalization by [programming Z3](#), in particular [Python API](#).

Is the order of the optimization functions important?

Types of optimization (in Z3)

Single-criteria optimization:

$OMT(\mathcal{LIR}\mathcal{A}\cup\mathcal{T})$, $OMT(\mathcal{BV}\cup\mathcal{T})$, $OMT(\mathcal{PB}\cup\mathcal{T})$ and MAXSMT solving [19].

Types of optimization (in Z3)

Single-criteria optimization:

$OMT(\mathcal{LIRAUT})$, $OMT(\mathcal{BVUT})$, $OMT(\mathcal{PBUT})$ and MAXSMT solving [19].

Multi-criteria optimization

To the best of our knowledge, in the SMT solver Z3, there are three ways to combine objective functions.

Types of optimization (in Z3)

Single-criteria optimization:

$OMT(\text{LIRAUT})$, $OMT(\text{BVUT})$, $OMT(\text{PBUT})$ and MAXSMT solving [19].

Multi-criteria optimization

To the best of our knowledge, in the SMT solver Z3, there are three ways to combine objective functions.

1. lexicographic combinations (by default) variant-int.smt2

Algorithm 3 Sequential algorithm for general objectives

- 1: **for** $t = 1$ to n **do**
- 2: Solve the single-objective problem:

$$\begin{aligned} & \max \quad f_t(x) \\ & \text{subject to} \quad x \in X, \\ & \quad f_k(x) \geq z_k \text{ for all } k \in 1, \dots, t-1. \end{aligned}$$

- 3: **if** the problem is infeasible or unbounded **then**
- 4: print "no solution"
- 5: **else**
- 6: Add as additional constraints the values of the decision variables x and $f_k(x) = z_t$
- 7: **end if**
- 8: **end for**

Types of optimization (in Z3) (cont'd)

2. **Boxes** are used to specify independent optima subject to given constraints:
variant-int-box.smt2

Types of optimization (in Z3) (cont'd)

2. **Boxes** are used to specify independent optima subject to given constraints:
variant-int-box.smt2
3. **Pareto optimization** involves more than one objective function to be optimized **simultaneously**. variant-int-pareto.smt2

Programming Z3 (Python API)

- ▶ variant-int.py

Programming Z3 (Python API)

- ▶ variant-int.py
- ▶ variant-bool.py

Feedback Part 1

Please fill out the form!



Contents

Motivation

Part 1: Optimization Modulo Theory - background and examples

Part 2: Optimization Modulo Theory - Case Study

Problem Specification

Problem Formalization

Model-driven approach: Formulation of the Satisfiability/Optimization Modulo Theory Problem

Data-driven approach: Graph Neural Network Formulation

Solution

Dataset generation

Training a GNN model for edge classification

Integrated GNN and Exact Techniques: Experimental Results

Future Work

Motivation

Model-Based and Data-Driven Design Methods

- Model-Based design methods
- Data-driven design methods
- Combination of model and data-based design methods
 - Incorporating learning in model-based methods
 - Incorporating models in AI/ML/CI methods
 - Reinforcement Learning and Approximate Dynamic Programming (ADP)
 - Adaptive Control and Learning Control
 -

• Frank Lewis and Derong Liu, Editors, *Reinforcement Learning and Approximate Dynamic Programming for Feedback Control*, Wiley, 2013.
• M Raissi, P Perdikaris, GE Karniadakis, "Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations," *Journal of Computational physics*, 2019
• J. Farrell and M. Polycarpou, *Adaptive Approximation Based Control: Unifying Neural, Fuzzy and Traditional Adaptive Approximation Approaches*, Wiley, 2006.

s.ucy.ac.cy

- ▶ The **data-driven** approach focuses on using data to drive the development and improvement of AI systems.
- ▶ The **model-based** approach focuses on developing a mathematical model of the system or process being studied.

Motivation

From Marios M. Polycarpou talk

The slide has a red header and a red decorative graphic of three spheres on the right. The main content is a bulleted list of design methods.

Model-Based and Data-Driven Design Methods

- Model-Based design methods
- Data-driven design methods
- Combination of model and data-based design methods
 - Incorporating learning in model-based methods
 - Incorporating models in AI/ML/CI methods
 - Reinforcement Learning and Approximate Dynamic Programming (ADP)
 - Adaptive Control and Learning Control
 -

• Frank Lewis and Derong Liu, Editors, *Reinforcement Learning and Approximate Dynamic Programming for Feedback Control*, Wiley, 2013.
• M Raisi, P Perdikaris, GE Karniadakis, "Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations," *Journal of Computational physics*, 2019.
• J. Farrell and M. Polycarpou, *Adaptive Approximation Based Control: Unifying Neural, Fuzzy and Traditional Adaptive Approximation Approaches*, Wiley, 2006.

slucy.ac.cy 22

- ▶ The **data-driven** approach focuses on using data to drive the development and improvement of AI systems.
- ▶ The **model-based** approach focuses on developing a mathematical model of the system or process being studied.

Motivation

Model-Based and Data-Driven Design Methods

- Model-Based design methods
- Data-driven design methods
- Combination of model and data-based design methods
 - Incorporating learning in model-based methods
 - Incorporating models in AI/ML/CI methods
 - Reinforcement Learning and Approximate Dynamic Programming (ADP)
 - Adaptive Control and Learning Control
 -

• Frank Lewis and Derong Liu, Editors, *Reinforcement Learning and Approximate Dynamic Programming for Feedback Control*, Wiley, 2013.
• M Raissi, P Perdikaris, GE Karniadakis, "Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations," *Journal of Computational physics*, 2019
• J. Farrell and M. Polycarpou, *Adaptive Approximation Based Control: Unifying Neural, Fuzzy and Traditional Adaptive Approximation Approaches*, Wiley, 2006.

s.ucy.ac.cy

- ▶ The **data-driven** approach focuses on using data to drive the development and improvement of AI systems. ↗ **graph neural networks**
- ▶ The **model-based** approach focuses on developing a mathematical model of the system or process being studied.

Motivation

Model-Based and Data-Driven Design Methods

- Model-Based design methods
- Data-driven design methods
- Combination of model and data-based design methods
 - Incorporating learning in model-based methods
 - Incorporating models in AI/ML/CI methods
 - Reinforcement Learning and Approximate Dynamic Programming (ADP)
 - Adaptive Control and Learning Control
 -

• Frank Lewis and Derong Liu, Editors, *Reinforcement Learning and Approximate Dynamic Programming for Feedback Control*, Wiley, 2013.
• M Ramezani, P Perdikaris, GE Karniadakis, "Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations," *Journal of Computational physics*, 2019
• J. Farrell and M. Polycarpou, *Adaptive Approximation Based Control: Unifying Neural, Fuzzy and Traditional Adaptive Approximation Approaches*, Wiley, 2006.

s.sucy.ac.cy

22

- ▶ The **data-driven** approach focuses on using data to drive the development and improvement of AI systems. \leadsto graph neural networks
- ▶ The **model-based** approach focuses on developing a mathematical model of the system or process being studied. \leadsto satisfiability/optimization modulo theory

Motivation (cont'd)

Advent of Cloud computing \rightsquigarrow loosely-coupled architecture \rightsquigarrow DevOps paradigm \rightsquigarrow application modeling \rightsquigarrow optimal deployment

Motivation (cont'd)

Advent of Cloud computing \rightsquigarrow loosely-coupled architecture \rightsquigarrow DevOps paradigm \rightsquigarrow application modeling \rightsquigarrow optimal deployment

Benefits of optimal deployment:

1. the synthesis of deployment plans that are **optimal by design**

Motivation (cont'd)

Advent of Cloud computing \rightsquigarrow loosely-coupled architecture \rightsquigarrow DevOps paradigm \rightsquigarrow application modeling \rightsquigarrow optimal deployment

Benefits of optimal deployment:

1. the synthesis of deployment plans that are **optimal by design**
2. the integration of such deployment plans into the application modeling process **enables formal reasoning** on a model of the deployed application.

Motivation (cont'd)

Advent of Cloud computing \rightsquigarrow loosely-coupled architecture \rightsquigarrow DevOps paradigm \rightsquigarrow application modeling \rightsquigarrow optimal deployment

Benefits of optimal deployment:

1. the synthesis of deployment plans that are **optimal by design**
2. the integration of such deployment plans into the application modeling process **enables formal reasoning** on a model of the deployed application.

Automated deployment of component-based applications in the Cloud consists of:

Motivation (cont'd)

Advent of Cloud computing \rightsquigarrow loosely-coupled architecture \rightsquigarrow DevOps paradigm \rightsquigarrow application modeling \rightsquigarrow optimal deployment

Benefits of optimal deployment:

1. the synthesis of deployment plans that are **optimal by design**
2. the integration of such deployment plans into the application modeling process **enables formal reasoning** on a model of the deployed application.

Automated deployment of component-based applications in the Cloud consists of:

1. **selection** of the computing resources,

Motivation (cont'd)

Advent of Cloud computing \rightsquigarrow loosely-coupled architecture \rightsquigarrow DevOps paradigm \rightsquigarrow application modeling \rightsquigarrow optimal deployment

Benefits of optimal deployment:

1. the synthesis of deployment plans that are **optimal by design**
2. the integration of such deployment plans into the application modeling process **enables formal reasoning** on a model of the deployed application.

Automated deployment of component-based applications in the Cloud consists of:

1. **selection** of the computing resources,
2. **distribution/assignment** of the application components over the available computing resources,

Motivation (cont'd)

Advent of Cloud computing \rightsquigarrow loosely-coupled architecture \rightsquigarrow DevOps paradigm \rightsquigarrow application modeling \rightsquigarrow optimal deployment

Benefits of optimal deployment:

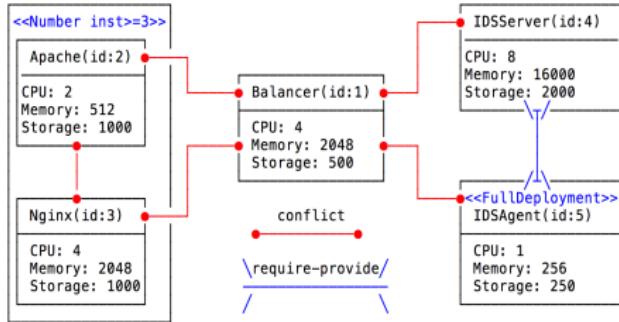
1. the synthesis of deployment plans that are **optimal by design**
2. the integration of such deployment plans into the application modeling process **enables formal reasoning** on a model of the deployed application.

Automated deployment of component-based applications in the Cloud consists of:

1. **selection** of the computing resources,
2. **distribution/assignment** of the application components over the available computing resources,
3. its **dynamic modification** to cope with peaks of user requests.

Motivating Example

Problem: finding the best offer for a Secure Web Container

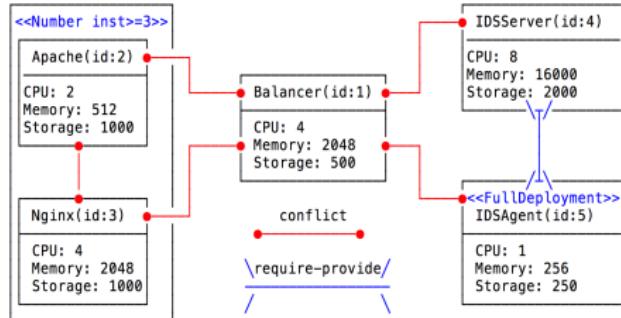


Components

- ▶ two Web Containers (e.g. Apache Tomcat or Nginx)
- ▶ a Balancer
- ▶ an IDSServer (Intrusion Detection System)
- ▶ an IDS Agent

Motivating Example

Problem: finding the best offer for a Secure Web Container



Components

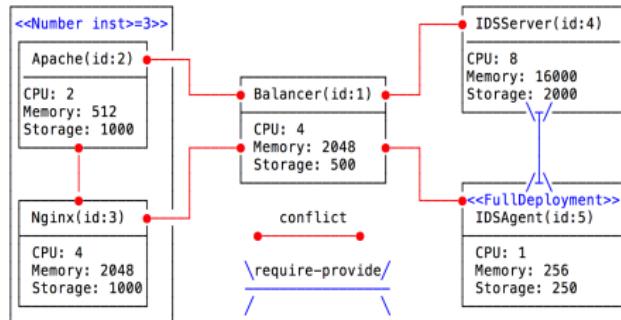
- ▶ two Web Containers (e.g. Apache Tomcat or Nginx)
- ▶ a Balancer
- ▶ an IDSServer (Intrusion Detection System)
- ▶ an IDS Agent

Constraints

- ▶ **Conflicts:** Apache and Nginx cannot be deployed on the same VM.
- ▶ **Conflicts:** Balancer needs exclusive use of machines.
- ▶ **Equal bound:** exactly one Balancer has to be instantiated.
- ▶ **Lower bound:** at least 3 instances of Apache and/or Nginx are required.
- ▶ **Require-provides:** one IDSServer for 10 IDS Agents.
- ▶ **Full deployment:** one instance of the IDS Agent on all VMs except for those containing the IDSServer and the Balancer.
- ▶ **Hardware constraints:** components hardware requirements.

Motivating Example

Problem: finding the best offer for a Secure Web Container



Components

- ▶ two Web Containers (e.g. Apache Tomcat or Nginx)
- ▶ a Balancer
- ▶ an IDSServer (Intrusion Detection System)
- ▶ an IDS Agent

Constraints

- ▶ **Conflicts**: Apache and Nginx cannot be deployed on the same VM.
- ▶ **Conflicts**: Balancer needs exclusive use of machines.
- ▶ **Equal bound**: exactly one Balancer has to be instantiated.
- ▶ **Lower bound**: at least 3 instances of Apache and/or Nginx are required.
- ▶ **Require-provides**: one IDSServer for 10 IDS Agents.
- ▶ **Full deployment**: one instance of the IDS Agent on all VMs except for those containing the IDSServer and the Balancer.
- ▶ **Hardware constraints**: components hardware requirements.

Goal: find a set of virtual machines (VMs) which satisfies the components requirements and lead to the minimum cost.

Motivating Example (cont'd)

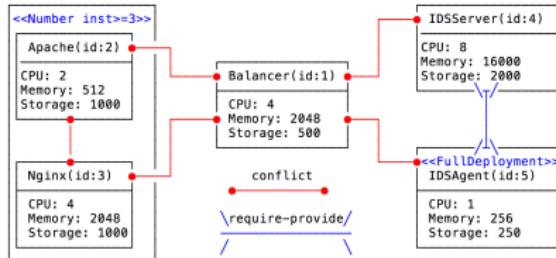
Spot Instance Prices

Spot Instances	Defined Duration for Linux	Defined Duration for Windows
Region: EU (Ireland)		
	Linux/UNIX Usage	Windows Usage
General Purpose - Current Generation		
t2.micro	\$0.0038 per Hour	\$0.0084 per Hour
t2.small	\$0.0075 per Hour	\$0.0165 per Hour
t2.medium	\$0.015 per Hour	\$0.033 per Hour
t2.large	\$0.0302 per Hour	\$0.0582 per Hour
t2.xlarge	\$0.0605 per Hour	\$0.1015 per Hour
t2.2xlarge	\$0.121 per Hour	\$0.183 per Hour
m3.medium	\$0.0073 per Hour	\$0.0633 per Hour
m3.large	\$0.0306 per Hour	\$0.1226 per Hour
m3.xlarge	\$0.0612 per Hour	\$0.2452 per Hour

Model	vCPU	CPU Credits / hour	Mem (GiB)	Storage
t2.nano	1	3	0.5	EBS-Only
t2.micro	1	6	1	EBS-Only
t2.small	1	12	2	EBS-Only
t2.medium	2	24	4	EBS-Only
t2.large	2	36	8	EBS-Only
t2.xlarge	4	54	16	EBS-Only
t2.2xlarge	8	81	32	EBS-Only

Remark: [snapshot from <https://aws.amazon.com/ec2/>] tens of thousands of price offers corresponding to different configurations and zones

Motivating Example (cont'd)



Example solution

- ▶ VM₁ (CPU:8, RAM: 15 GB, Storage: 2000 GB, Price: 0.0526 \$/hour):
Nginx + IDS Agent
- ▶ VM₂ (CPU:4, RAM: 7.5 GB, Storage: 2000 GB, Price: 0.0283 \$/hour):
Balancer
- ▶ VM₃ (CPU:4, RAM: 30 GB, Storage: 2000 GB, Price: 0.0644 \$/hour):
IDSServer
- ▶ VM₄ (CPU:4, RAM: 7.5 GB, Storage: 2000 GB, Price: 0.0283 \$/hour):
Apache + IDS Agent
- ▶ VM₅ (CPU:4, RAM: 7.5 GB, Storage: 2000 GB, Price: 0.0283 \$/hour):
Apache + IDS Agent

Contents

Motivation

Part 1: Optimization Modulo Theory - background and examples

Part 2: Optimization Modulo Theory - Case Study

Problem Specification

Problem Formalization

Model-driven approach: Formulation of the Satisfiability/Optimization Modulo Theory Problem

Data-driven approach: Graph Neural Network Formulation

Solution

Dataset generation

Training a GNN model for edge classification

Integrated GNN and Exact Techniques: Experimental Results

Future Work

Problem Specification

Automated deployment of component-based applications in the Cloud consists of:

1. **selection** of the computing resources,
2. **distribution/assignment** of the application components over the available computing resources,
3. **dynamic modification** to cope with peaks of user requests.

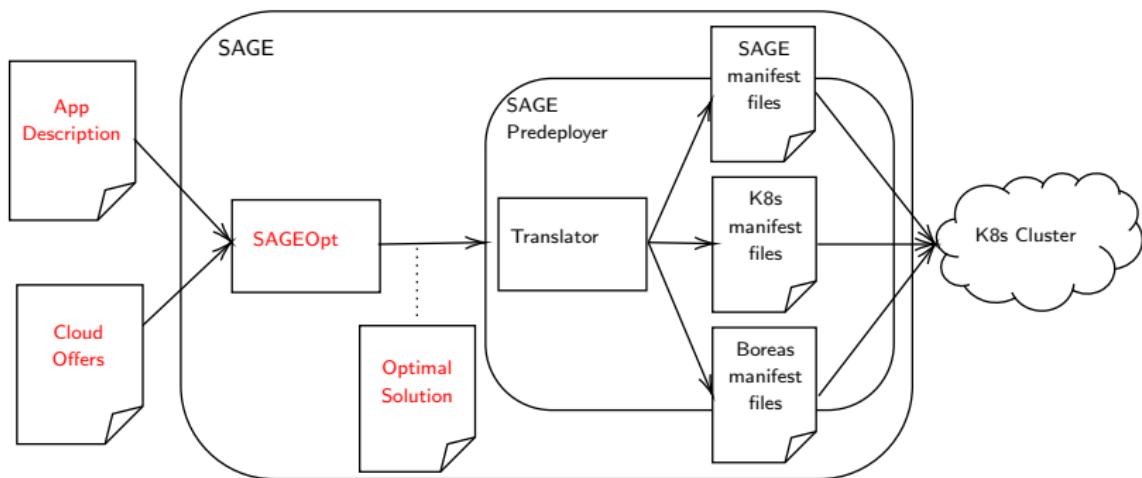


Figure: SAGE General Architecture

Contents

Motivation

Part 1: Optimization Modulo Theory - background and examples

Part 2: Optimization Modulo Theory - Case Study

Problem Specification

Problem Formalization

Model-driven approach: Formulation of the Satisfiability/Optimization Modulo Theory Problem

Data-driven approach: Graph Neural Network Formulation

Solution

Dataset generation

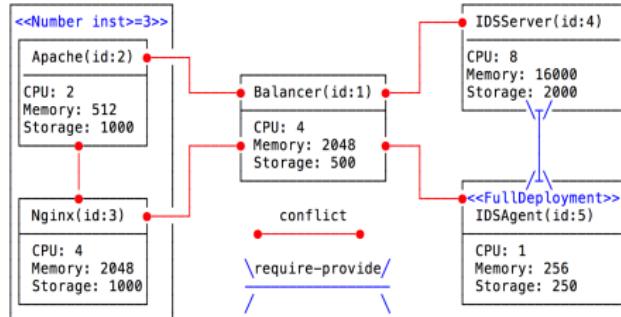
Training a GNN model for edge classification

Integrated GNN and Exact Techniques: Experimental Results

Future Work

Recall Motivating Example

Problem: finding the best offer for a Secure Web Container



Components

- ▶ two Web Containers (e.g. Apache Tomcat or Nginx)
- ▶ a Balancer
- ▶ an IDSServer (Intrusion Detection System)
- ▶ an IDS Agent

Constraints

- ▶ **Conflicts**: Apache and Nginx cannot be deployed on the same VM.
- ▶ **Conflicts**: Balancer needs exclusive use of machines.
- ▶ **Equal bound**: exactly one Balancer has to be instantiated.
- ▶ **Lower bound**: at least 3 instances of Apache and/or Nginx are required.
- ▶ **Require-provides**: one IDSServer for 10 IDS Agents.
- ▶ **Full deployment**: one instance of the IDS Agent on all VMs except for those containing the IDSServer and the Balancer.
- ▶ **Hardware constraints**: components hardware requirements.

Goal: find a set of virtual machines (VMs) which satisfies the components requirements and lead to the minimum cost.

Model-driven approach: Formulation of the Satisfiability/Optimization Modulo Theory Problem

General constraints

$$\text{Basic allocation} \quad \sum_{k=1}^M a_{ik} \geq 1 \quad \forall i = \overline{1, N}$$

$$\text{Occupancy} \quad \sum_{i=1}^N a_{ik} \geq 1 \Rightarrow v_k = 1 \quad \forall k = \overline{1, M}$$

$$\text{Capacity} \quad \sum_{i=1}^N a_{ik} \cdot R_i^h \leq F_{t_k}^h \quad \forall k = \overline{1, M}, \forall h = \overline{1, H}$$

$$\text{Link} \quad v_k = 1 \wedge t_k = o \Rightarrow \bigwedge_{h=1}^H (r_k^h = F_{t_k}^h) \wedge p_k = P_{t_k} \quad \forall o = \overline{1, O}, O \in \mathbb{N}^*$$

$$\sum_{i=1}^N a_{ik} = 0 \Rightarrow t_k = 0 \quad \forall k = \overline{1, M}$$

where:

- ▶ $R_i^h \in \mathbb{N}^*$ is the hardware requirement of type h of the component i ;
- ▶ $F_{t_k}^h \in \mathbb{N}^*$ is the hardware characteristic h of the VM of type t_k .

Problem Formalization (cont'd)

Application-specific constraints

$$\text{Conflicts} \quad a_{ik} + a_{jk} \leq 1$$

$$\forall k = \overline{1, M}, \forall (i, j) \quad R_{ij} = 1$$

$$\text{Co-location} \quad a_{ik} = a_{jk}$$

$$\forall k = \overline{1, M}, \forall (i, j) \quad D_{ij} = 1$$

$$\text{Exclusive deployment}$$

$$H\left(\sum_{k=1}^M a_{i_1 k}\right) + \dots + H\left(\sum_{k=1}^M a_{i_q k}\right) = 1 \quad \text{for fixed } q \in \{1, \dots, N\}$$

$$H(u) = \begin{cases} 1 & u > 0 \\ 0 & u = 0 \end{cases}$$

Require- Provide

$$n_{ij} \sum_{k=1}^M a_{ik} \leq m_{ij} \sum_{k=1}^M a_{jk} \quad \forall (i, j) Q_{ij}(n_{ij}, m_{ij}) = 1$$

$$0 \leq n \sum_{k=1}^M a_{jk} - \sum_{k=1}^M a_{ik} < n \quad n, n_{ij}, m_{ij} \in \mathbb{N}^*$$

where:

- $R_{ij} = 1$ if components i and j are in conflict (can not be placed in the same VM);
- $D_{ij} = 1$ if components i and j must be co-located (must be placed in the same VM);
- $Q_{ij}(n, m) = 1$ if C_i requires at least n instances of C_j and C_j can serve at most m instances of C_i

Problem Formalization (cont'd)

Application-specific constraints

$$\text{Full deployment} \quad \sum_{k=1}^M \left(a_{ik} + \mathcal{H} \left(\sum_{j, \mathcal{R}_{ij}=1} a_{jk} \right) \right) = \sum_{k=1}^M v_k$$

Deployment with bounded number of instances

$$\sum_{i \in \bar{C}} \sum_{k=1}^M a_{ik} \langle \text{op} \rangle n \quad |\bar{C}| \leq N, \langle \text{op} \rangle \in \{=, \leq, \geq\}, n \in \mathbb{N}$$

Find:

- ▶ assignment matrix a with binary entries $a_{ik} \in \{0, 1\}$ for $i = \overline{1, N}$, $k = \overline{1, M}$, which are interpreted as follows:

$$a_{ik} = \begin{cases} 1 & \text{if } C_i \text{ is assigned to } V_k \\ 0 & \text{if } C_i \text{ is not assigned to } V_k. \end{cases}$$

- ▶ the type selection vector t with integer entries t_k for $k = \overline{1, M}$, representing the type (from a predefined set) of each VM leased.

Such that: the leasing price is minimal $\sum_{k=1}^M v_k \cdot p_k$

Graph Neural Network Formulation

The first step in solving the **edge classification problem** is to model it as **graph data**.

Graph Neural Network Formulation

The first step in solving the **edge classification problem** is to model it as **graph data**.

Heterogeneous graph:

Graph Neural Network Formulation

The first step in solving the **edge classification problem** is to model it as **graph data**.

Heterogeneous graph:

- ▶ component nodes

- ▶ VM nodes

Graph Neural Network Formulation

The first step in solving the **edge classification problem** is to model it as **graph data**.

Heterogeneous graph:

- ▶ **component nodes**
 - ▶ The **features** are (*ID, CPU, Mem, Sto, FullDepl, UpperB, LowerB, EqualB*).
- ▶ **VM nodes**

Graph Neural Network Formulation

The first step in solving the **edge classification problem** is to model it as **graph data**.

Heterogeneous graph:

- ▶ **component nodes**
 - ▶ The **features** are (*ID, CPU, Mem, Sto, FullDepl, UpperB, LowerB, EqualB*).
 - ▶ **Edges** between component nodes are determined by the application-specific constraints except *FullDepl, UpperB, LowerB, EqualB*.
- ▶ **VM nodes**

Graph Neural Network Formulation

The first step in solving the **edge classification problem** is to model it as **graph data**.

Heterogeneous graph:

- ▶ **component nodes**
 - ▶ The **features** are (*ID, CPU, Mem, Sto, FullDepl, UpperB, LowerB, EqualB*).
 - ▶ **Edges** between component nodes are determined by the application-specific constraints except *FullDepl, UpperB, LowerB, EqualB*.
 - ▶ Edges are specified using **one-hot encoding** on the possible application-specific constraints: (*Conflict, Co-location, RequireProvide, ExclusiveDeployment, UpperB, LowerB, EqualB*).
- ▶ **VM nodes**

Graph Neural Network Formulation

The first step in solving the **edge classification problem** is to model it as **graph data**.

Heterogeneous graph:

- ▶ **component nodes**
 - ▶ The **features** are (*ID, CPU, Mem, Sto, FullDepl, UpperB, LowerB, EqualB*).
 - ▶ **Edges** between component nodes are determined by the application-specific constraints except *FullDepl, UpperB, LowerB, EqualB*.
 - ▶ Edges are specified using **one-hot encoding** on the possible application-specific constraints: (*Conflict, Co-location, RequireProvide, ExclusiveDeployment, UpperB, LowerB, EqualB*).
- ▶ **VM nodes**
 - ▶ The **features** are (*CPU, Mem, Sto, Price*)

Graph Neural Network Formulation

The first step in solving the **edge classification problem** is to model it as **graph data**.

Heterogeneous graph:

- ▶ **component nodes**
 - ▶ The **features** are (*ID, CPU, Mem, Sto, FullDepl, UpperB, LowerB, EqualB*).
 - ▶ **Edges** between component nodes are determined by the application-specific constraints except *FullDepl, UpperB, LowerB, EqualB*.
 - ▶ Edges are specified using **one-hot encoding** on the possible application-specific constraints: (*Conflict, Co-location, RequireProvide, ExclusiveDeployment, UpperB, LowerB, EqualB*).
- ▶ **VM nodes**
 - ▶ The **features** are (*CPU, Mem, Sto, Price*)
 - ▶ **No edges** between two VM nodes!

Graph Neural Network Formulation

The first step in solving the **edge classification problem** is to model it as **graph data**.

Heterogeneous graph:

- ▶ **component nodes**
 - ▶ The **features** are (*ID, CPU, Mem, Sto, FullDepl, UpperB, LowerB, EqualB*).
 - ▶ **Edges** between component nodes are determined by the application-specific constraints except *FullDepl, UpperB, LowerB, EqualB*.
 - ▶ Edges are specified using **one-hot encoding** on the possible application-specific constraints: (*Conflict, Co-location, RequireProvide, ExclusiveDeployment, UpperB, LowerB, EqualB*).
- ▶ **VM nodes**
 - ▶ The **features** are (*CPU, Mem, Sto, Price*)
 - ▶ **No edges** between two VM nodes!

Feature scaling:

Graph Neural Network Formulation

The first step in solving the **edge classification problem** is to model it as **graph data**.

Heterogeneous graph:

- ▶ **component nodes**
 - ▶ The **features** are (*ID, CPU, Mem, Sto, FullDepl, UpperB, LowerB, EqualB*).
 - ▶ **Edges** between component nodes are determined by the application-specific constraints except *FullDepl, UpperB, LowerB, EqualB*.
 - ▶ Edges are specified using **one-hot encoding** on the possible application-specific constraints: (*Conflict, Co-location, RequireProvide, ExclusiveDeployment, UpperB, LowerB, EqualB*).
- ▶ **VM nodes**
 - ▶ The **features** are (*CPU, Mem, Sto, Price*)
 - ▶ **No edges** between two VM nodes!

Feature scaling:

- ▶ to [0, 1] is needed in order to prevent one feature from dominating while preserving min-max relationships, crucial for prediction algorithms
- ▶ is needed for the hardware specifications/requirements and the VM price.

Graph Neural Network Formulation

The first step in solving the **edge classification problem** is to model it as **graph data**.

Heterogeneous graph:

- ▶ **component nodes**
 - ▶ The **features** are (*ID, CPU, Mem, Sto, FullDepl, UpperB, LowerB, EqualB*).
 - ▶ **Edges** between component nodes are determined by the application-specific constraints except *FullDepl, UpperB, LowerB, EqualB*.
 - ▶ Edges are specified using **one-hot encoding** on the possible application-specific constraints: (*Conflict, Co-location, RequireProvide, ExclusiveDeployment, UpperB, LowerB, EqualB*).
- ▶ **VM nodes**
 - ▶ The **features** are (*CPU, Mem, Sto, Price*)
 - ▶ **No edges** between two VM nodes!

Feature scaling:

- ▶ to [0, 1] is needed in order to prevent one feature from dominating while preserving min-max relationships, crucial for prediction algorithms
- ▶ is needed for the hardware specifications/requirements and the VM price.

Edge classification

Graph Neural Network Formulation

The first step in solving the **edge classification problem** is to model it as **graph data**.

Heterogeneous graph:

- ▶ **component nodes**
 - ▶ The **features** are (*ID, CPU, Mem, Sto, FullDepl, UpperB, LowerB, EqualB*).
 - ▶ **Edges** between component nodes are determined by the application-specific constraints except *FullDepl, UpperB, LowerB, EqualB*.
 - ▶ Edges are specified using **one-hot encoding** on the possible application-specific constraints: (*Conflict, Co-location, RequireProvide, ExclusiveDeployment, UpperB, LowerB, EqualB*).
- ▶ **VM nodes**
 - ▶ The **features** are (*CPU, Mem, Sto, Price*)
 - ▶ **No edges** between two VM nodes!

Feature scaling:

- ▶ to [0, 1] is needed in order to prevent one feature from dominating while preserving min-max relationships, crucial for prediction algorithms
- ▶ is needed for the hardware specifications/requirements and the VM price.

Edge classification

- ▶ Initially, all edges between a node of type component and one of type VM are of type *unlinked*.

Graph Neural Network Formulation

The first step in solving the **edge classification problem** is to model it as **graph data**.

Heterogeneous graph:

- ▶ **component nodes**
 - ▶ The **features** are (*ID, CPU, Mem, Sto, FullDepl, UpperB, LowerB, EqualB*).
 - ▶ **Edges** between component nodes are determined by the application-specific constraints except *FullDepl, UpperB, LowerB, EqualB*.
 - ▶ Edges are specified using **one-hot encoding** on the possible application-specific constraints: (*Conflict, Co-location, RequireProvide, ExclusiveDeployment, UpperB, LowerB, EqualB*).
- ▶ **VM nodes**
 - ▶ The **features** are (*CPU, Mem, Sto, Price*)
 - ▶ **No edges** between two VM nodes!

Feature scaling:

- ▶ to [0, 1] is needed in order to prevent one feature from dominating while preserving min-max relationships, crucial for prediction algorithms
- ▶ is needed for the hardware specifications/requirements and the VM price.

Edge classification

- ▶ Initially, all edges between a node of type component and one of type VM are of type *unlinked*.

The task is to implement a GNN model in order to predict the type (linked/unlinked) for all the edges.

Contents

Motivation

Part 1: Optimization Modulo Theory - background and examples

Part 2: Optimization Modulo Theory - Case Study

Problem Specification

Problem Formalization

Model-driven approach: Formulation of the Satisfiability/Optimization Modulo Theory Problem

Data-driven approach: Graph Neural Network Formulation

Solution

Dataset generation

Training a GNN model for edge classification

Integrated GNN and Exact Techniques: Experimental Results

Future Work

Solution

1. Generate the dataset which is used to train a GNN model of the application to be deployed. This dataset, representing optimal deployment plans, is obtained by multiple runs of the exact solver previously developed by us.

Solution

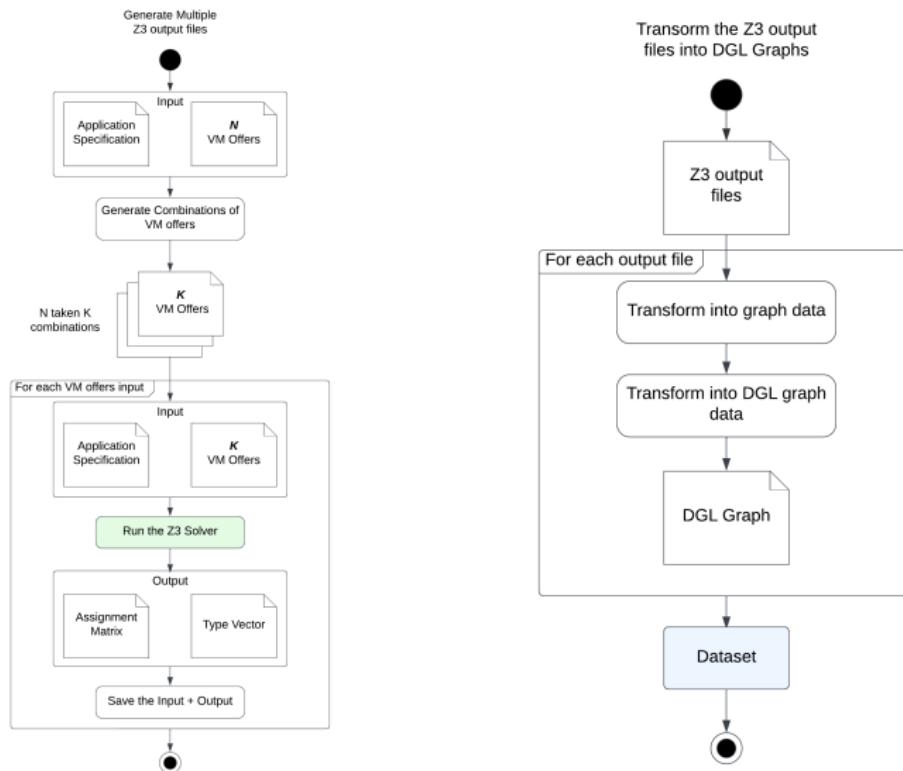
1. Generate the dataset which is used to train a GNN model of the application to be deployed. This dataset, representing optimal deployment plans, is obtained by multiple runs of the exact solver previously developed by us.
2. Train a GNN model which *predicts* the assignments of components to VMs as well as the VM Offers.

Solution

1. Generate the dataset which is used to train a GNN model of the application to be deployed. This dataset, representing optimal deployment plans, is obtained by multiple runs of the exact solver previously developed by us.
2. Train a GNN model which *predicts* the assignments of components to VMs as well as the VM Offers.
3. Transform the predictions into *soft constraints* to guide the search exploration of the Base solver towards an optimal solution.

Dataset generation

Large dataset to train the model $\binom{20}{15} \approx 15000$ different VM Offers inputs.



Training a GNN model for edge classification

Supervised GNN learning approach:

Training a GNN model for edge classification

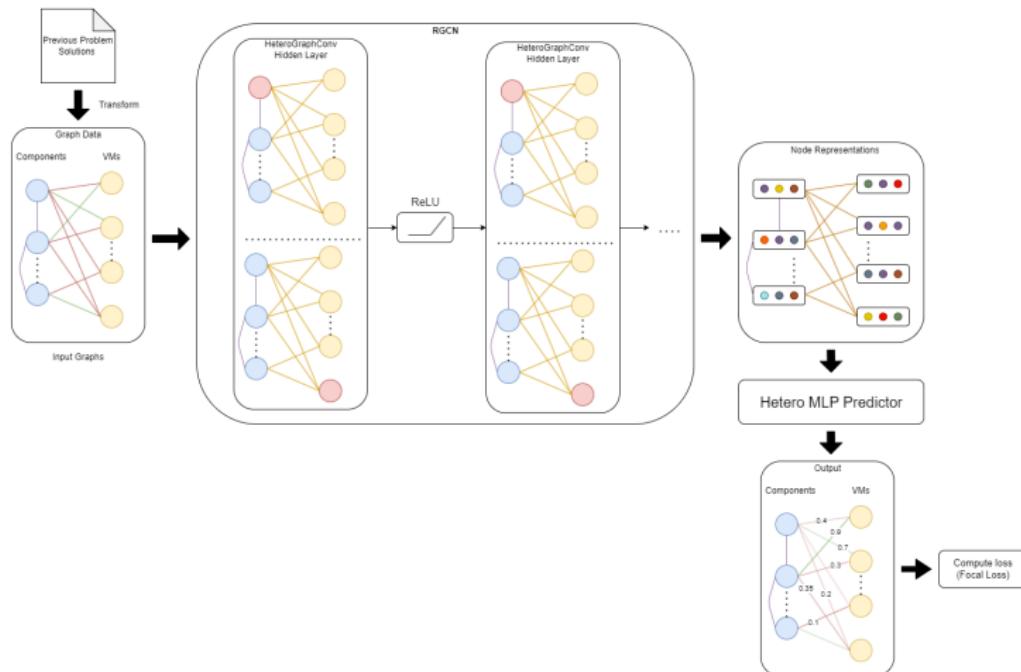
Supervised GNN learning approach:

1. **Data Preparation**: graph representation and nodes and edges feature extraction.
2. **Graph Construction**: application graph's structure (nodes, edges, and their relationships)

Training a GNN model for edge classification

Supervised GNN learning approach:

3. Choosing Model Architecture which allows heterogeneity modeling and edge classification.



Training a GNN model for edge classification

Supervised GNN learning approach:

4. The **loss function** is focal loss suitable for imbalanced sets.

Training a GNN model for edge classification

Supervised GNN learning approach:

4. The **loss function** is focal loss suitable for imbalanced sets.
5. **Adam optimizer** with the default parameters provided by DGL library.

Training a GNN model for edge classification

Supervised GNN learning approach:

4. The **loss function** is focal loss suitable for imbalanced sets.
5. **Adam optimizer** with the default parameters provided by DGL library.
6. **Model Training, Validation and Testing**

Training a GNN model for edge classification

Supervised GNN learning approach:

4. The **loss function** is focal loss suitable for imbalanced sets.
5. **Adam optimizer** with the default parameters provided by DGL library.
6. **Model Training, Validation and Testing**
 - ▶ train, test, validation: 60%, 20%, 20%

Training a GNN model for edge classification

Supervised GNN learning approach:

4. The **loss function** is focal loss suitable for imbalanced sets.
5. **Adam optimizer** with the default parameters provided by DGL library.
6. **Model Training, Validation and Testing**
 - ▶ train, test, validation: 60%, 20%, 20%
 - ▶ batch size = 50

Training a GNN model for edge classification

Supervised GNN learning approach:

4. The **loss function** is focal loss suitable for imbalanced sets.
5. **Adam optimizer** with the default parameters provided by DGL library.
6. **Model Training, Validation and Testing**
 - ▶ train, test, validation: 60%, 20%, 20%
 - ▶ batch size = 50
 - ▶ training dataset size: min = 50, max = 10000

Training a GNN model for edge classification

Supervised GNN learning approach:

4. The **loss function** is focal loss suitable for imbalanced sets.
5. **Adam optimizer** with the default parameters provided by DGL library.
6. **Model Training, Validation and Testing**
 - ▶ train, test, validation: 60%, 20%, 20%
 - ▶ batch size = 50
 - ▶ training dataset size: min = 50, max = 10000
 - ▶ #epochs: min = 30, max = 600

Training a GNN model for edge classification

Supervised GNN learning approach:

4. The **loss function** is focal loss suitable for imbalanced sets.
5. **Adam optimizer** with the default parameters provided by DGL library.
6. **Model Training, Validation and Testing**
 - ▶ train, test, validation: 60%, 20%, 20%
 - ▶ batch size = 50
 - ▶ training dataset size: min = 50, max = 10000
 - ▶ #epochs: min = 30, max = 600

#	Sample Size	#Ep	Acc	Time	Pred. T Links	Pred. F Links	GT True Links
3	50	200	0.95	21.36	7	10	8
7	100	100	0.95	21.92	7	13	
11	100	400	0.95	87.92	8	13	

Integrated GNN and Exact Techniques: Experimental Results

- ▶ the scalability of the GNN approach for increasing number of VM offers, possibly previously unseen (see table), and

Integrated GNN and Exact Techniques: Experimental Results

- ▶ the scalability of the GNN approach for increasing number of VM offers, possibly previously unseen (see table), and
- ▶ the generalization of the GNN approach for applications characterized by similar constraints between components but with different hardware requirements (see the paper [9]).

Experimental Analysis (cont'd)

Explanation of the FV symmetry breaker

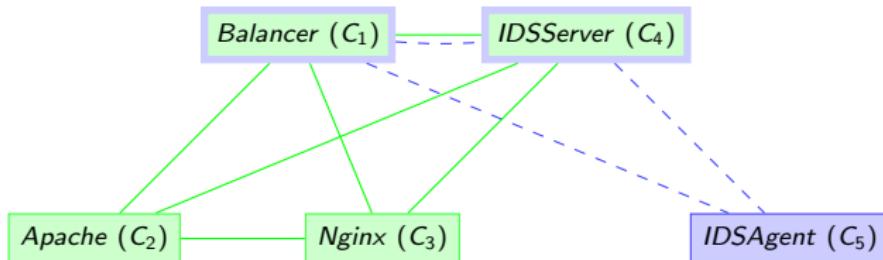


Figure: Secure Web Container conflict graph. The components with green background belong to the clique \overline{G} .

	VM_1	VM_2	VM_3	VM_4	VM_5	VM_6
C_1	1	0	0	0	0	0
C_2	0	0	1	0	1	0
C_3	0	0	0	1	0	0
C_4	0	1	0	0	0	0
C_5	0	0	1	1	1	0

Table: Effect of FV symmetry breaking strategy

Integrated GNN and Exact Techniques: Experimental Results

- ▶ the scalability of the GNN approach for increasing number of VM offers, possibly previously unseen (see table). Base = Z3

#o	Solver	Model#3	Model#7	Model#11	Opt. Price	
20	Base	0.24				
	Base+FVPR	0.11				
	Base+GNN	0.12	0.12	0.07	3.759	
	Base+FVPR+GNN	0.12	0.09	0.10		
40	Base	0.54			2.676	
	Base+FVPR	0.27				
	Base+GNN	0.28	0.33	0.28		
	Base+FVPR+GNN	0.29	0.28	0.29		
250	Base	2.82			1.622	
	Base+FVPR	0.98				
	Base+GNN	0.76	0.77	1		
	Base+FVPR+GNN	1.40	1.39	1.15		
500	Base	8.71			1.582	
	Base+FVPR	2.42				
	Base+GNN	4.56	2.92	1.5		
	Base+FVPR+GNN	3.09	3	3.01		
27	Base	0.26			2.400	
	Base+FVPR	0.09				
	Base+GNN	0.10	0.14	0.14		
	Base+FVPR+GNN	0.10	0.10	0.07		

Experimental Analysis (cont'd)

Using the GNN Prediction as Soft Constraints

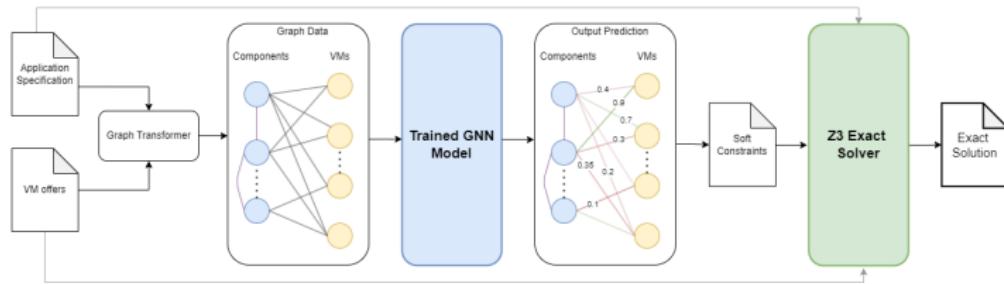
Formalization of the assignment predictions as a binary 3D tensor pred with $\text{pred}_{iko} \in \{0, 1\}$ for $i = \overline{1, N}$, $k = \overline{1, M}$ and $o = \overline{1, O}$:

$$\text{pred}_{iko} = \begin{cases} 1 & \text{if } C_i \text{ is assigned to } V_j \text{ of type } O_o \\ 0 & \text{if } C_i \text{ is not assigned to } V_j \text{ of type } O_o \end{cases}$$

From tensors to soft constraints

$$\exists o \in \overline{1, O} \text{ s.t. } \text{pred}_{iko} = 1 \implies a_{ik} = 1 \wedge \bigwedge_{h=1}^H (r_k^h = F_o^h) \wedge p_k = P_o$$
$$\nexists o \in \overline{1, O} \text{ s.t. } \text{pred}_{iko} = 1 \implies a_{ik} = 0$$

Diagram describing the integration of the GNN model with the Base solver



Experimental Analysis (cont'd)

Using the GNN Prediction as Soft Constraints The *pred* tensor for the *first component* of the Secure Web Container application, generated from running the GNN model prediction on the case study application with 10 VM offers, looks like:

$$\begin{pmatrix} 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{pmatrix}$$

where $M = 6$ rows and $O = 10$ columns.

The corresponding soft constraints are:

- ▶ for assignment matrix a :

```
(assert-soft (= a11 0))      (assert-soft (= a13 1))      (assert-soft (= a15 0))
(assert-soft (= a12 1))      (assert-soft (= a14 1))      (assert-soft (= a16 0))
```

- ▶ for type vector t :

```
(assert-soft (and (= PriceProv2 8.403)...))
(assert-soft (and (= PriceProv3 8.403)...))
(assert-soft (and (= PriceProv4 0.093)...))
```

where the predictions obtained for the VM offers type were $t_2 = 7$, $t_3 = 7$, $t_4 = 5$. For example, the VM Offer 5 has the specification $(1, 3750, 1000, 0.093)$.

Experimental Analysis (cont'd)

Implementation of soft constraints in Z3 optimization

Optimization in Z3:

1. If the soft constraints are added before the optimization function, then the solver tries to satisfy as many as possible and will use the intermediate results for the optimization function (multi-criterial optimization with lexicographic option by default). \leadsto hence the soft constraints might destroy the actual optimim

Experimental Analysis (cont'd)

Implementation of soft constraints in Z3 optimization

Optimization in Z3:

1. If the soft constraints are added before the optimization function, then the solver tries to satisfy as many as possible and will use the intermediate results for the optimization function (multi-criterial optimization with lexicographic option by default). \rightsquigarrow hence the soft constraints might destroy the actual optimim
2. If the soft constraints are declared after the optimization function, the optimal value will be found without the soft constraints \rightsquigarrow this does not exploit the benefits of soft constraints in guiding the solver to find faster the solution.

Experimental Analysis (cont'd)

Implementation of soft constraints in Z3 optimization

Optimization in Z3:

1. If the soft constraints are added before the optimization function, then the solver tries to satisfy as many as possible and will use the intermediate results for the optimization function (multi-criterial optimization with lexicographic option by default). \rightsquigarrow hence the soft constraints might destroy the actual optimim
2. If the soft constraints are declared after the optimization function, the optimal value will be found without the soft constraints \rightsquigarrow this does not exploit the benefits of soft constraints in guiding the solver to find faster the solution.

We simulated soft constraints with *pseudo-boolean and cardinality constraints*.

Experimental Analysis (cont'd)

Implementation of soft constraints in Z3 optimization

Optimization in Z3:

1. If the soft constraints are added before the optimization function, then the solver tries to satisfy as many as possible and will use the intermediate results for the optimization function (multi-criterial optimization with lexicographic option by default). \rightsquigarrow hence the soft constraints might destroy the actual optimim
2. If the soft constraints are declared after the optimization function, the optimal value will be found without the soft constraints \rightsquigarrow this does not exploit the benefits of soft constraints in guiding the solver to find faster the solution.

We simulated soft constraints with *pseudo-boolean and cardinality constraints*.

Pseudo-Boolean constraint

A linear *pseudo-Boolean constraint* has the form: $\sum_j a_i l_j \triangleright b$ where a_i and b are integer constants, l_j are literals and \triangleright is a relational operator.

Experimental Analysis (cont'd)

Implementation of soft constraints in Z3 optimization

Optimization in Z3:

1. If the soft constraints are added before the optimization function, then the solver tries to satisfy as many as possible and will use the intermediate results for the optimization function (multi-criterial optimization with lexicographic option by default). \rightsquigarrow hence the soft constraints might destroy the actual optimim
2. If the soft constraints are declared after the optimization function, the optimal value will be found without the soft constraints \rightsquigarrow this does not exploit the benefits of soft constraints in guiding the solver to find faster the solution.

We simulated soft constraints with *pseudo-boolean and cardinality constraints*.

Pseudo-Boolean constraint

A linear *pseudo-Boolean constraint* has the form: $\sum_j a_i l_j \triangleright b$ where a_i and b are integer constants, l_j are literals and \triangleright is a relational operator.

Cardinality Constraint

A *cardinality constraint* is a constraint on the number of literals which are true among a given set of literals.

Experimental Analysis (cont'd)

Implementation of soft constraints in Z3 optimization

Optimization in Z3:

1. If the soft constraints are added before the optimization function, then the solver tries to satisfy as many as possible and will use the intermediate results for the optimization function (multi-criterial optimization with lexicographic option by default). \rightsquigarrow hence the soft constraints might destroy the actual optimim
2. If the soft constraints are declared after the optimization function, the optimal value will be found without the soft constraints \rightsquigarrow this does not exploit the benefits of soft constraints in guiding the solver to find faster the solution.

We simulated soft constraints with *pseudo-boolean* and *cardinality constraints*.

Pseudo-Boolean constraint

A linear *pseudo-Boolean constraint* has the form: $\sum_j a_i l_j \triangleright b$ where a_i and b are integer constants, l_j are literals and \triangleright is a relational operator.

Cardinality Constraint

A *cardinality constraint* is a constraint on the number of literals which are true among a given set of literals.

In our case: `atmost(k, {x1, x2, ..., xn})` is true if and only if at most k literals among x_1, x_2, \dots, x_n are true.

Contents

Motivation

Part 1: Optimization Modulo Theory - background and examples

Part 2: Optimization Modulo Theory - Case Study

Problem Specification

Problem Formalization

Model-driven approach: Formulation of the Satisfiability/Optimization Modulo Theory Problem

Data-driven approach: Graph Neural Network Formulation

Solution

Dataset generation

Training a GNN model for edge classification

Integrated GNN and Exact Techniques: Experimental Results

Future Work

Future Work

- ▶ Investigate better the timings of Base+GNN and Base+FV+GNN on this use case and others.

Future Work

- ▶ Investigate better the timings of Base+GNN and Base+FV+GNN on this use case and others.
 - ▶ Investigate the characteristics of the datasets.
 - ▶ Improve the GNN model.
 - ▶ New case studies, from different application domains.

References I

-  IBM ILOG CPLEX Optimization Studio.
Version, 12(1987-2018):1, 1987.
-  R. F. Araujo, H. F. Albuquerque, I. V. De Bessa, L. C. Cordeiro, and J. E. Chaves Filho.
Counterexample guided inductive optimization based on satisfiability modulo theories.
Science of Computer Programming, 165:3–23, 2018.
-  H. Barbosa, C. Barrett, M. Brain, G. Kremer, H. Lachnitt, M. Mann, A. Mohamed, M. Mohamed, A. Niemetz, A. Nötzli, et al.
CVC5: A versatile and industrial-strength SMT solver.
In *International Conference on Tools and Algorithms for the Construction and Analysis of Systems*, pages 415–442. Springer, 2022.
-  N. Bjørner, A. Phan, and L. Fleckenstein.
 ν Z - an optimizing SMT solver.
In *Tools and Algorithms for the Construction and Analysis of Systems - 21st International Conference, TACAS 2015, London, UK, April 11-18, 2015. Proceedings*, pages 194–199, 2015.

References II

-  R. Bruttomesso, E. Pek, N. Sharygina, and A. Tsitovich.
The OpenSMT solver.
In *Tools and Algorithms for the Construction and Analysis of Systems: 16th International Conference, TACAS 2010, Held as Part of the Joint European Conferences on Theory and Practice of Software, ETAPS 2010, Paphos, Cyprus, March 20-28, 2010. Proceedings 16*, pages 150–153. Springer, 2010.
-  G. Chu, P. Stuckey, A. Schutt, T. Ehlers, G. Gange, and K. Francis.
Chuffed, a lazy clause generation solver.
-  L. de Moura and N. Bjørner.
Z3: An efficient SMT solver.
In C. R. Ramakrishnan and J. Rehof, editors, *Tools and Algorithms for the Construction and Analysis of Systems*, pages 337–340, Berlin, Heidelberg, 2008. Springer Berlin Heidelberg.
-  B. Dutertre.
Yices 2 manual.
Computer Science Laboratory, SRI International, Tech. Rep., 2014.

References III



M. Eraşcu.

Fast and exact synthesis of application deployment plans using graph neural networks and satisfiability modulo theory.

In *2024 International Joint Conference on Neural Networks (IJCNN)*, pages 1–10, 2024.



K. Fazekas, F. Bacchus, and A. Biere.

Implicit hitting set algorithms for maximum satisfiability modulo theories.

In *Automated Reasoning: 9th International Joint Conference, IJCAR 2018, Held as Part of the Federated Logic Conference, FloC 2018, Oxford, UK, July 14–17, 2018, Proceedings 9*, pages 134–151. Springer, 2018.



G. Kovásznai, C. Biró, and B. Erdélyi.

Puli – a problem-specific OMT solver.

In *Proc. 16th International Workshop on Satisfiability Modulo Theories (SMT 2018)*, volume 371, 2018.



D. Larraz, A. Oliveras, E. Rodríguez-Carbonell, and A. Rubio.

Minimal-model-guided approaches to solving polynomial constraints and extensions.

In *International Conference on Theory and Applications of Satisfiability Testing*, pages 333–350. Springer, 2014.

References IV

-  Y. Li, A. Albarghouthi, Z. Kincaid, A. Gurfinkel, and M. Chechik.
Symbolic optimization with SMT solvers.
In *The 41st Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages, POPL '14, San Diego, CA, USA, January 20-21, 2014*, pages 607–618, 2014.
-  A. Nadel and V. Ryvchin.
Bit-vector optimization.
In *Tools and Algorithms for the Construction and Analysis of Systems: 22nd International Conference, TACAS 2016, Held as Part of the European Joint Conferences on Theory and Practice of Software, ETAPS 2016, Eindhoven, The Netherlands, April 2-8, 2016, Proceedings* 22, pages 851–867. Springer, 2016.
-  N. Nethercote, P. J. Stuckey, R. Becket, S. Brand, G. J. Duck, and G. Tack.
MiniZinc: Towards a standard CP modelling language.
In *International Conference on Principles and Practice of Constraint Programming*, pages 529–543. Springer, 2007.
-  L. Perron and V. Furnon.
OR-Tools.

References V

-  C. Schulte, M. Lagerkvist, and G. Tack.
Gecode.
Software download and online material at the website:
<http://www.gecode.org>, pages 11–13, 2006.
-  R. Sebastiani and P. Trentin.
OptiMathSAT: A tool for optimization modulo theories.
In D. Kroening and C. S. Păsăreanu, editors, *Computer Aided Verification*,
pages 447–454, Cham, 2015. Springer International Publishing.
-  P. Trentin.
Optimization Modulo Theories with OptiMathSAT.
Phd thesis, University of Trento, May 2019.